

Differentiable NAS

Shusen Wang

<http://wangshusen.github.io/>

Reference

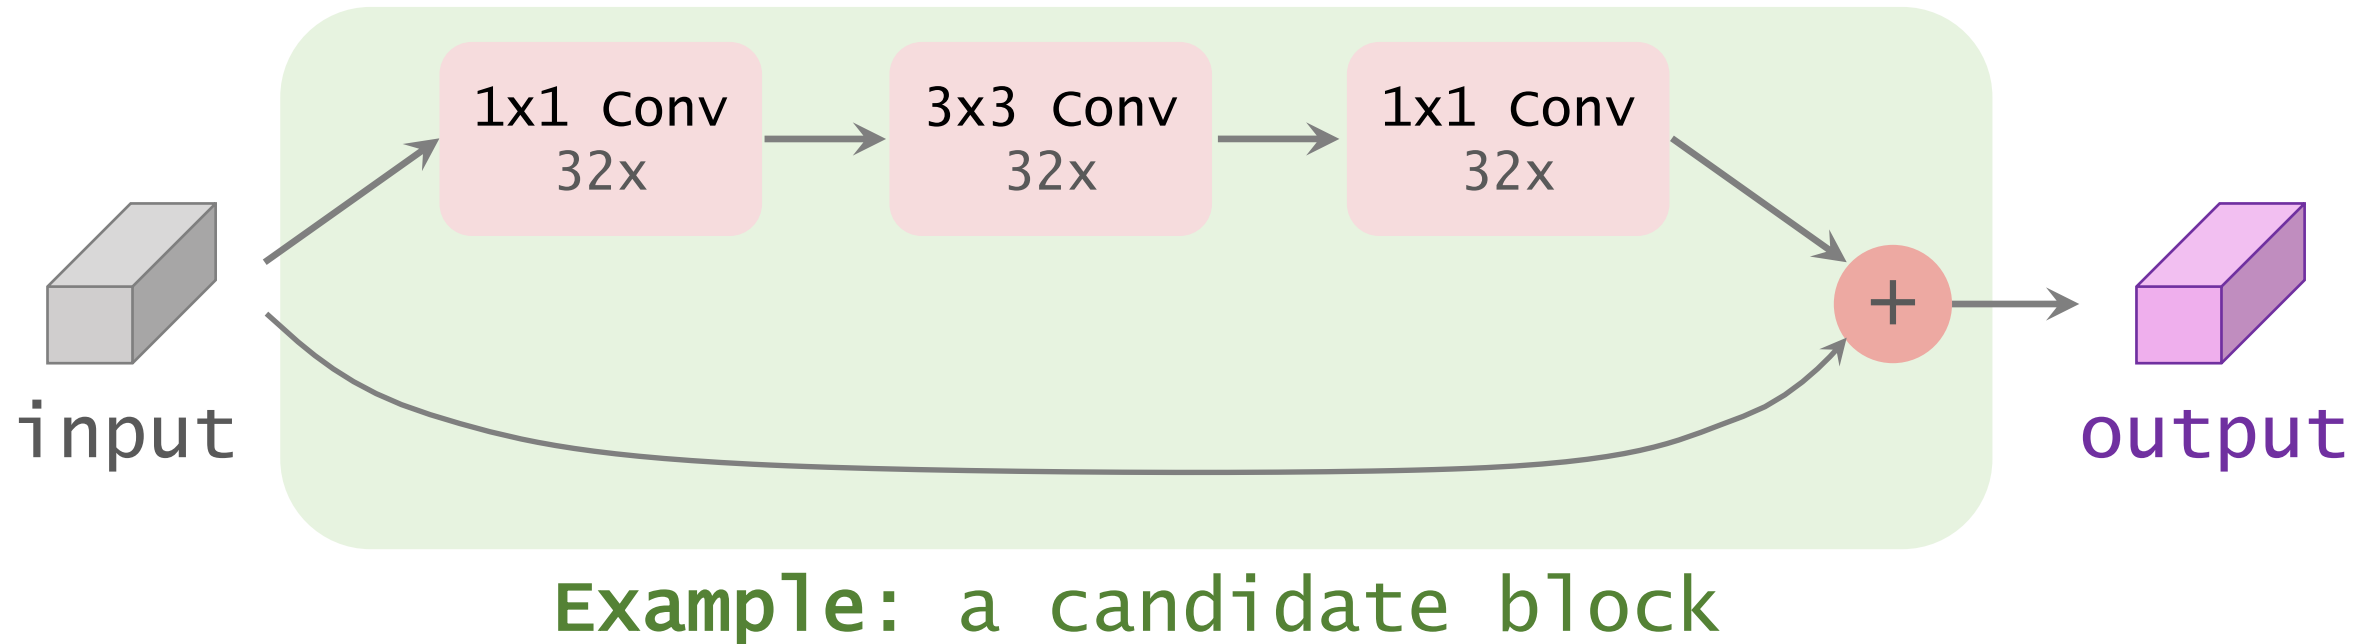
-  1. Liu, Simonyan, & Yang. DARTS: Differentiable Architecture Search. In *ICLR*, 2019.
-  2. Wu et al. FBNet: Hardware-Aware Efficient ConvNet Design via Differentiable Neural Architecture Search. In *CVPR*, 2019.

Basic Idea

- User manually defines some (e.g., 9) candidate blocks.

Basic Idea

- User manually defines some (e.g., 9) candidate blocks.




Basic Idea

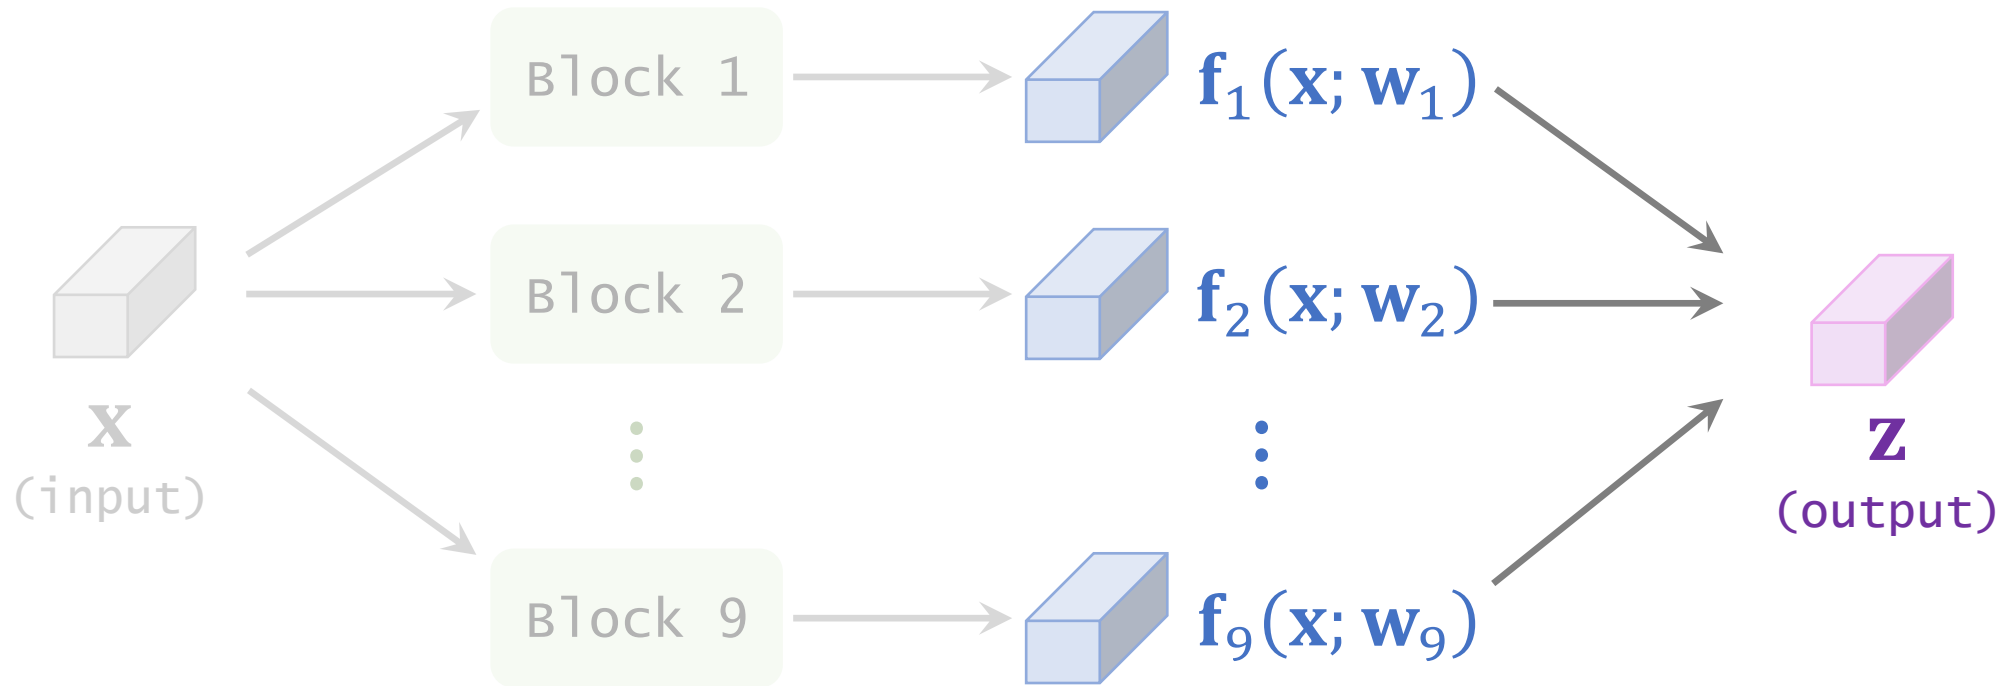
- User manually defines some (e.g., 9) candidate blocks.
- User specifies the number of layers, e.g., 20 layers.
- Each layer can be one of the 9 candidate blocks.
- Size of search space (i.e., # of possible architectures) is 9^{20} .

Super-net

One Layer of Super-net

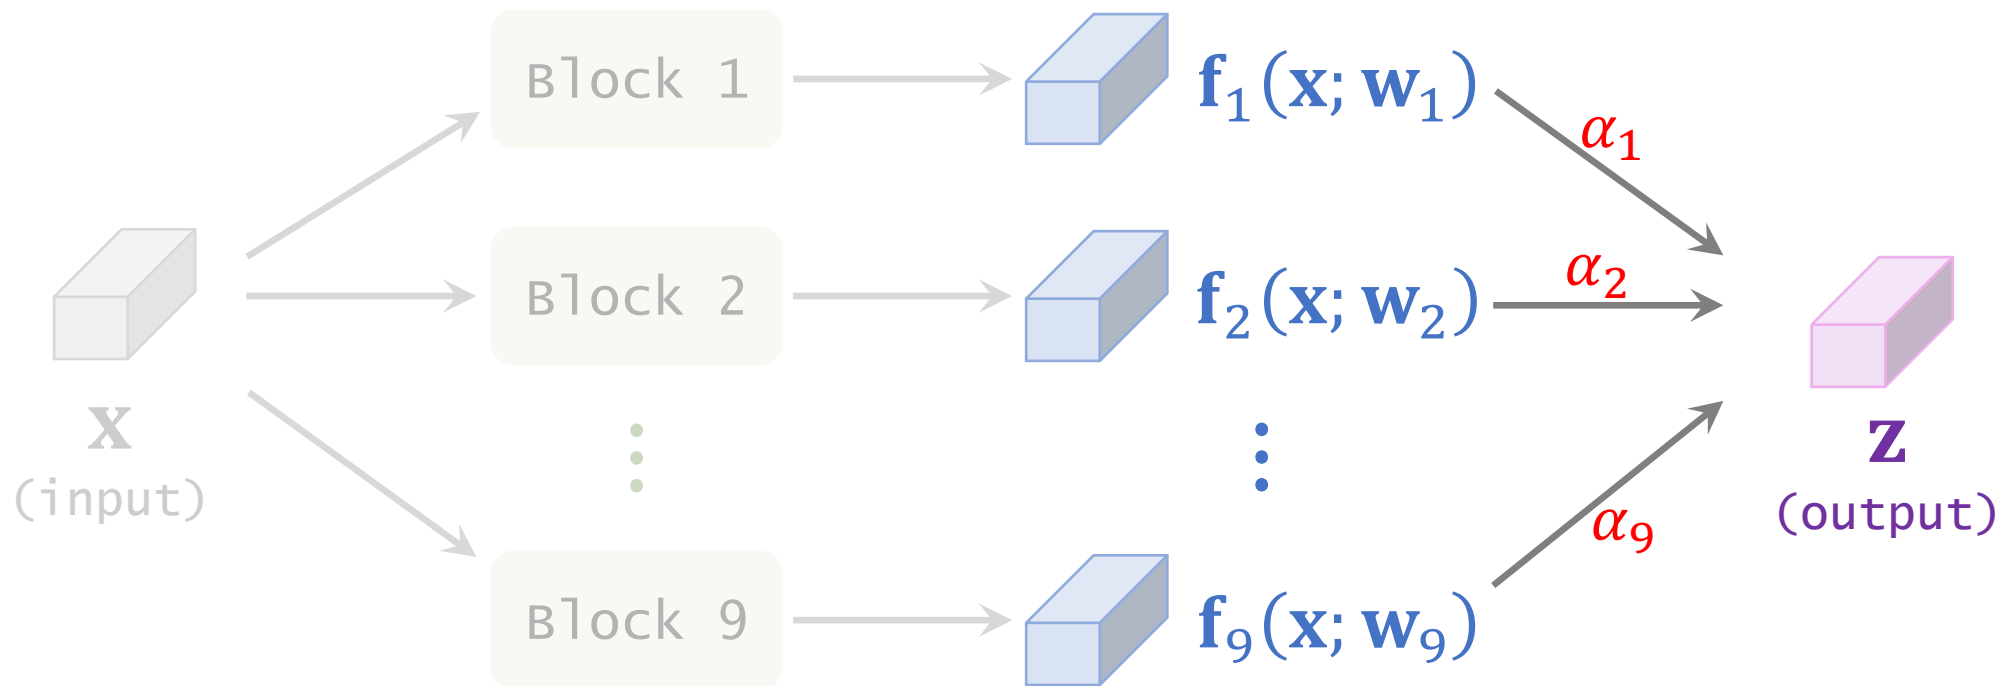


One Layer of Super-net



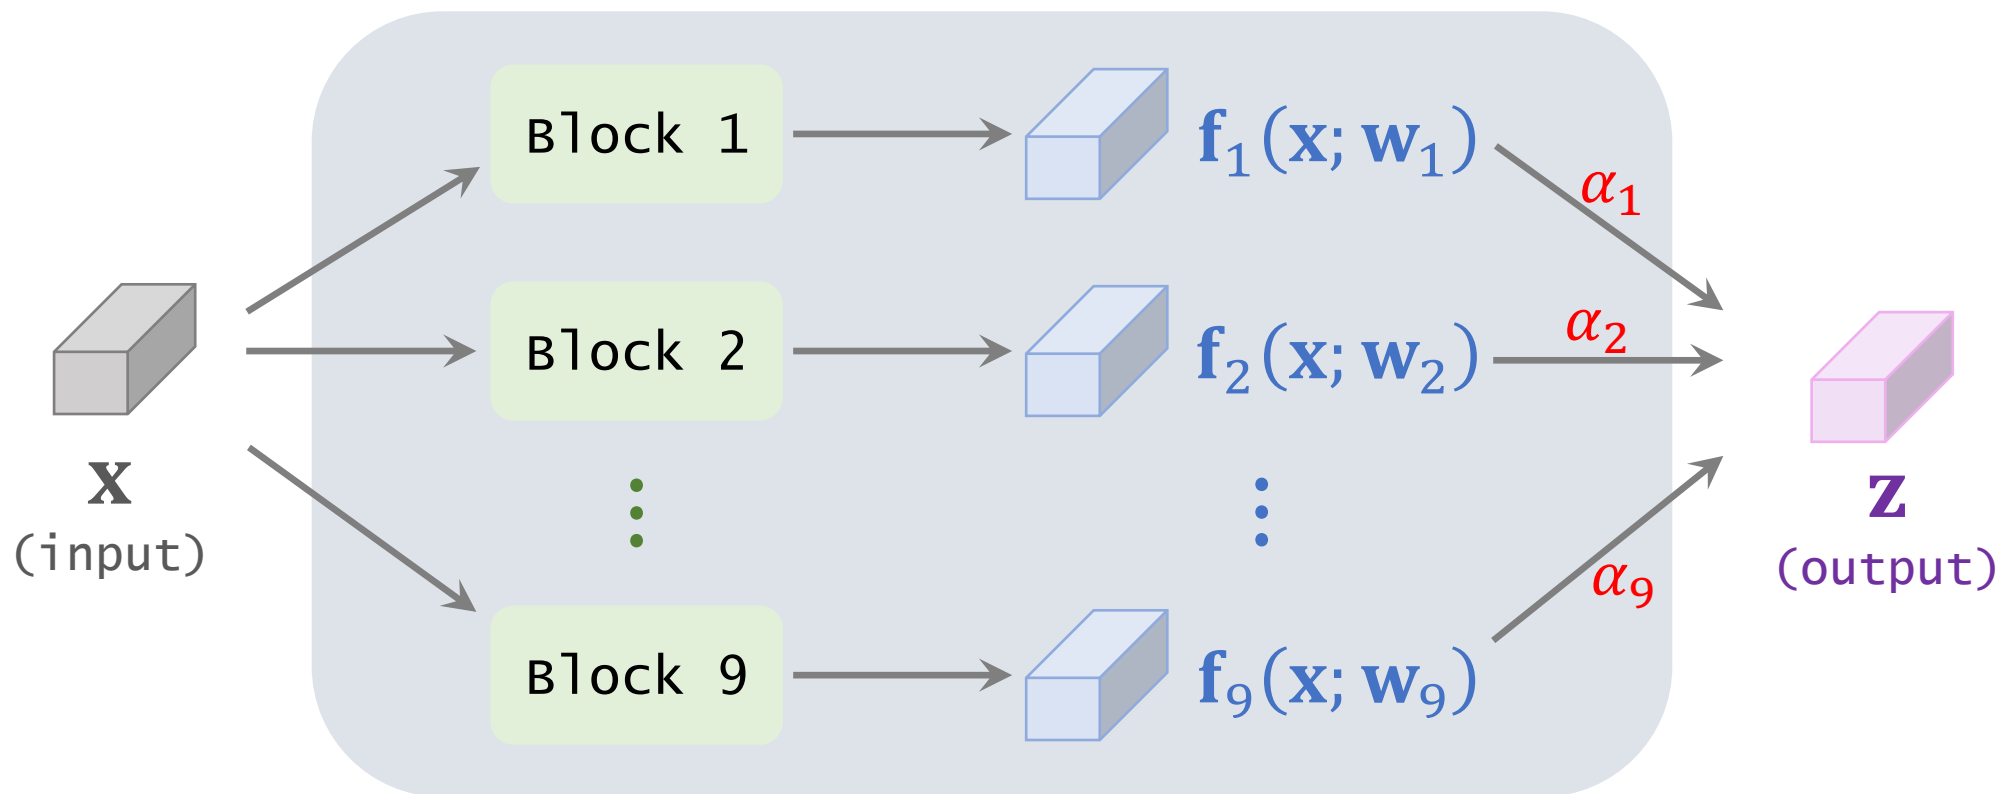
One Layer of Super-net

- $[\alpha_1, \dots, \alpha_9]$ = Softmax($\theta_1, \dots, \theta_9$).
- Output: $\mathbf{z} = \sum_{j=1}^9 \alpha_j \cdot \mathbf{f}_j(\mathbf{x}; \mathbf{w}_j)$.




One Layer of Super-net

One Layer




The super-net has 20 layers; each layer contains 9 parallel blocks.

Layer 1



...

Layer 20



Trainable Parameters of Super-net

- Each layer has the following trainable parameters:
 - $\mathbf{w}_1, \dots, \mathbf{w}_9$ (tensors): parameters of the 9 blocks.
 - $\theta_1, \dots, \theta_9$ (scalars): parameters that determine the weights, $\alpha_1, \dots, \alpha_9$.
- Layers do not share parameters.
 - Each layer has its own parameters, $\mathbf{w}_1, \dots, \mathbf{w}_9$ and $\theta_1, \dots, \theta_9$.
 - Parameters are not shared across layers.

Trainable Parameters of Super-net


- Blocks: $j = 1, \dots, 9$.
- Layers: $l = 1, \dots, 20$.
- Trainable parameters of the l -th layer and j -th block:
 - $\mathbf{w}_j^{(l)}$ (tensors) and $\theta_j^{(l)}$ (a scalar).
- All the trainable parameters of the super-net:
 - $\mathcal{W} = \{ \mathbf{w}_j^{(l)} \}_{j,l}$ and $\Theta = \{ \theta_j^{(l)} \}_{j,l}$.

Train the Super-net

- $\mathbf{x}_1, \dots, \mathbf{x}_n$: training images.
- $\mathbf{y}_1, \dots, \mathbf{y}_n$: targets (aka labels).
- $\mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)$: a prediction made by the 20-layer super-net.

Train the Super-net


- $\mathbf{x}_1, \dots, \mathbf{x}_n$: training images.
- $\mathbf{y}_1, \dots, \mathbf{y}_n$: targets (aka labels).
- $\mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)$: a prediction made by the 20-layer super-net.
- Learn \mathcal{W} and Θ from the training set by minimizing the cross-entropy loss:

$$\min_{\mathcal{W}, \Theta} \frac{1}{n} \sum_{i=1}^n \text{Loss}(\mathbf{y}_i, \mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)).$$


Train the Super-net


- $\mathbf{x}_1, \dots, \mathbf{x}_n$: training images.
- $\mathbf{y}_1, \dots, \mathbf{y}_n$: targets (aka labels).
- $\mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)$: a prediction made by the 20-layer super-net.
- Learn \mathcal{W} and Θ from the training set by minimizing the cross-entropy loss:

$$\min_{\mathcal{W}, \Theta} \frac{1}{n} \sum_{i=1}^n \text{Loss}(\mathbf{y}_i, \mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)).$$




Knowing the optimal Θ , we have the weights $\alpha_j^{(l)} = \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$.

Layer 1




...

Layer 20




Knowing the optimal Θ , we have the weights $\alpha_j^{(l)} = \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$.

Layer 1




...

Layer 20




Knowing the optimal Θ , we have the weights $\alpha_j^{(l)} = \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$.

Layer 1




...

Layer 20




For each layer, select the block that has the biggest weight, α .

Layer 1




...

Layer 20




For each layer, select the block that has the biggest weight, α .

Layer 1




Layer 20




For each layer, select the block that has the biggest weight, α .

Layer 1



...

Layer 20



Computational Efficient Design

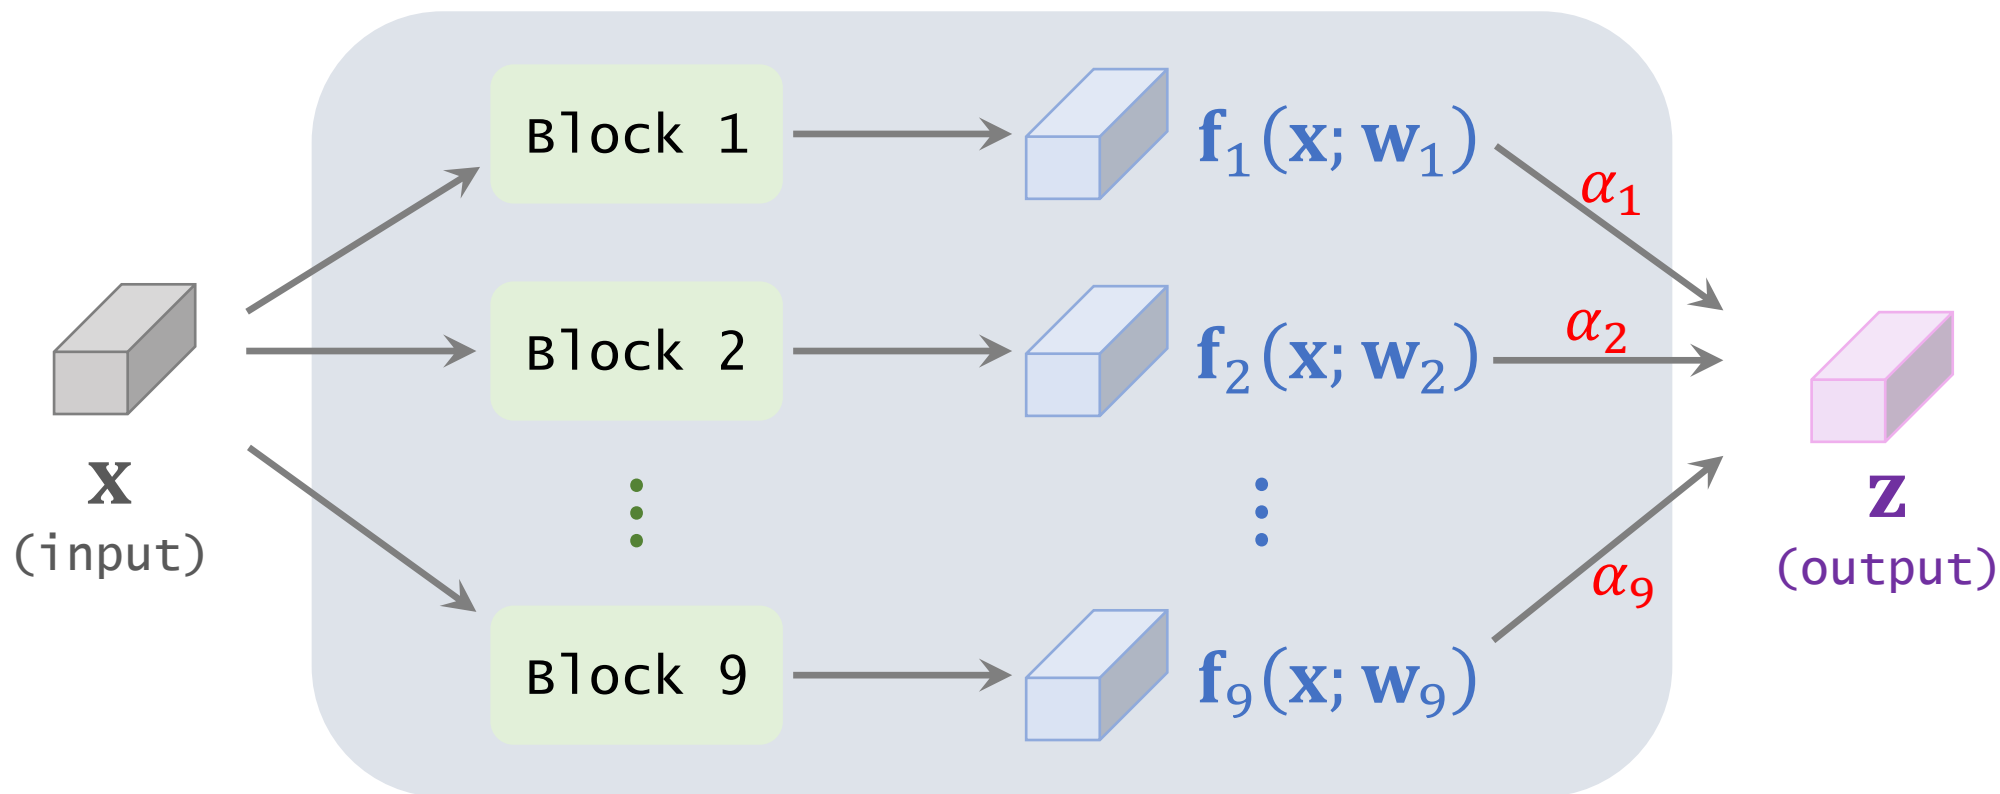
Reference:

1. Wu et al. [FBNet: Hardware-aware efficient convnet design via differentiable neural architecture search](#). In *CVPR*, 2019.

Basic Idea

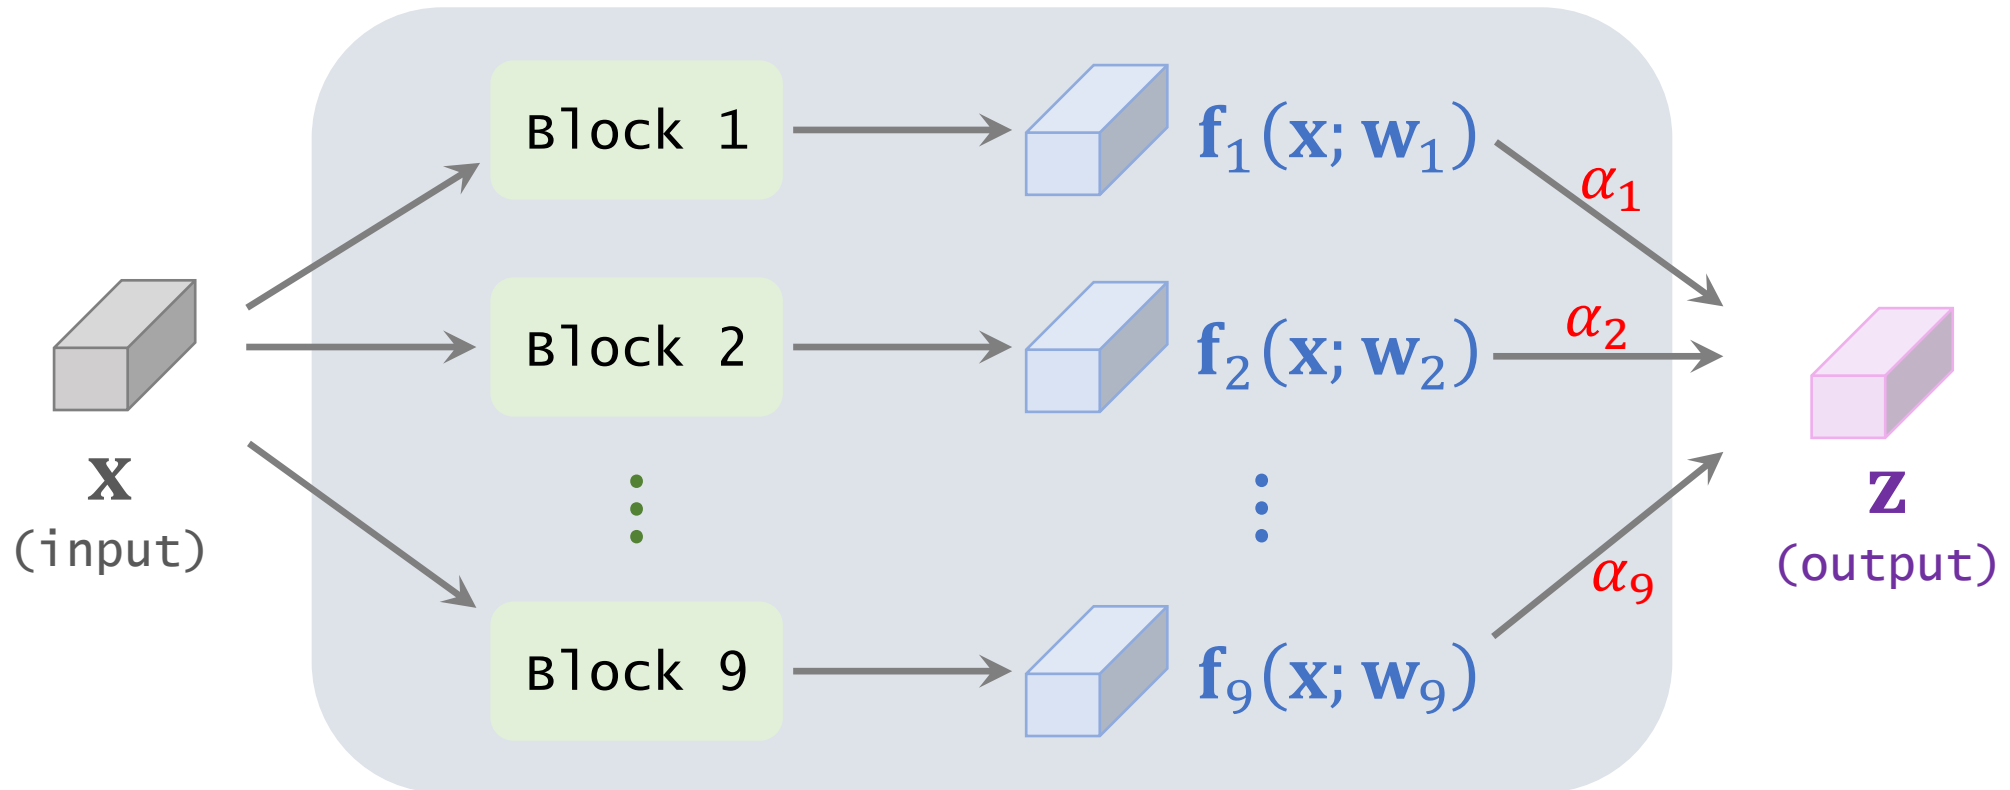
- A trained CNN takes as input an image and makes a prediction.
- Small latency (i.e., time cost of prediction making) is preferable.
- Latency can be considered during architecture search.
 - Different candidate blocks cause different accuracies and different latencies.
 - Trade off accuracy and latency.

One Layer of the Super-net



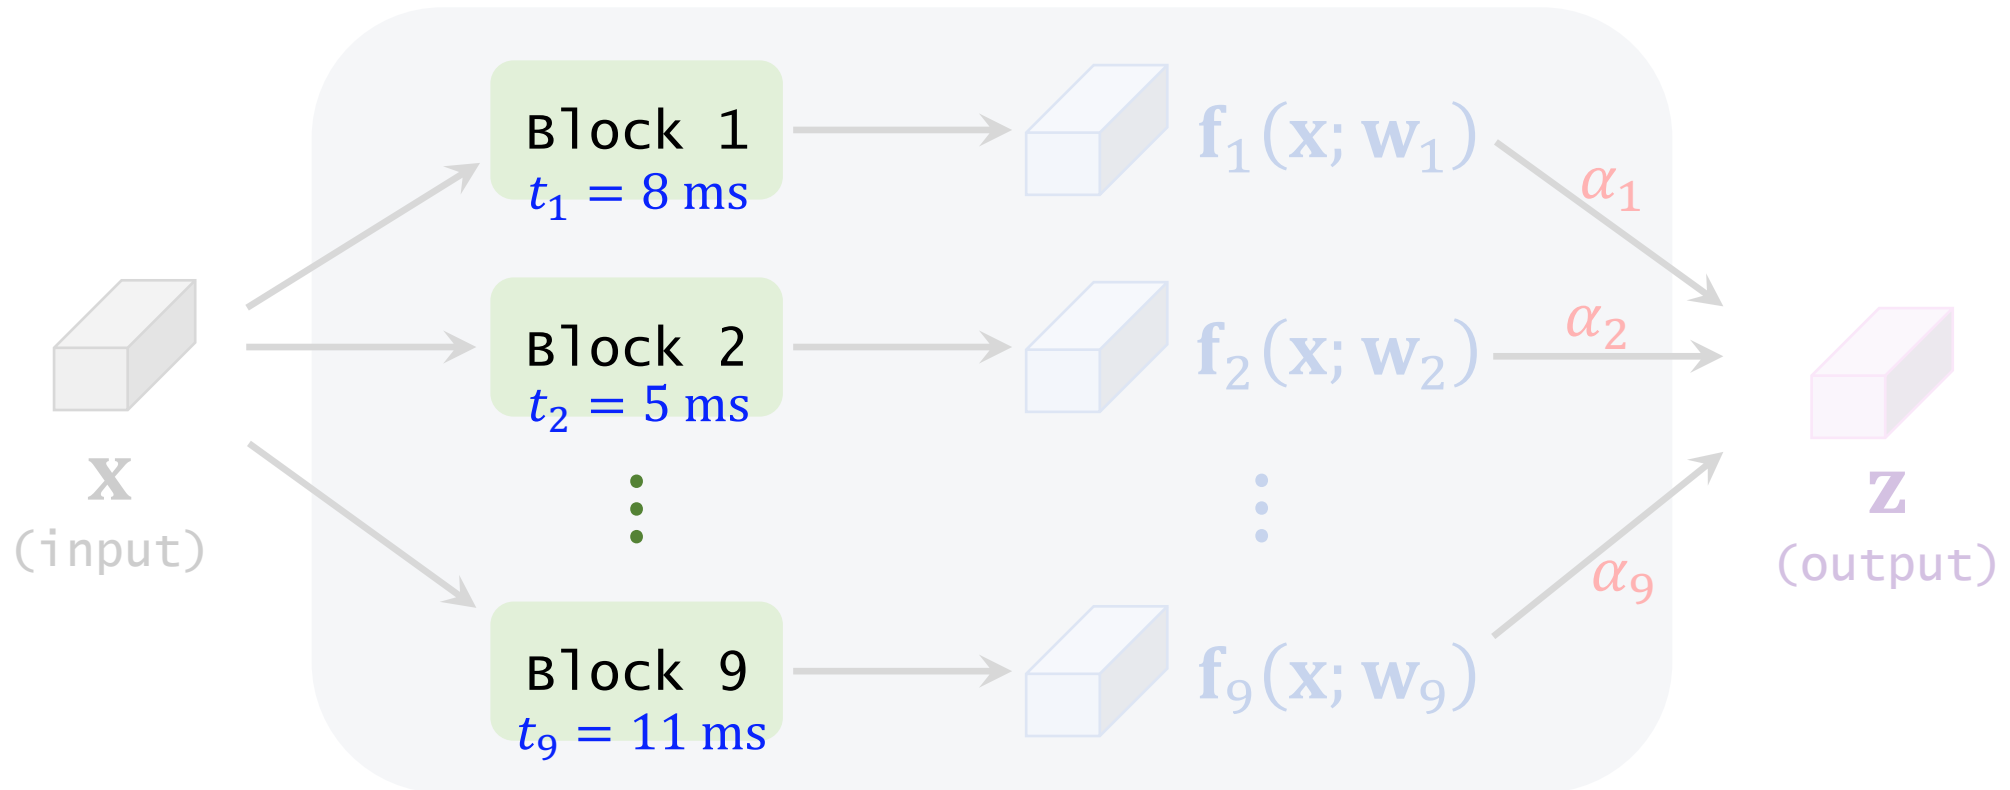
Latency

- Suppose the selected CNN will be deployed to iPhone 12.
- On iPhone 12, measure the latency caused by each block.



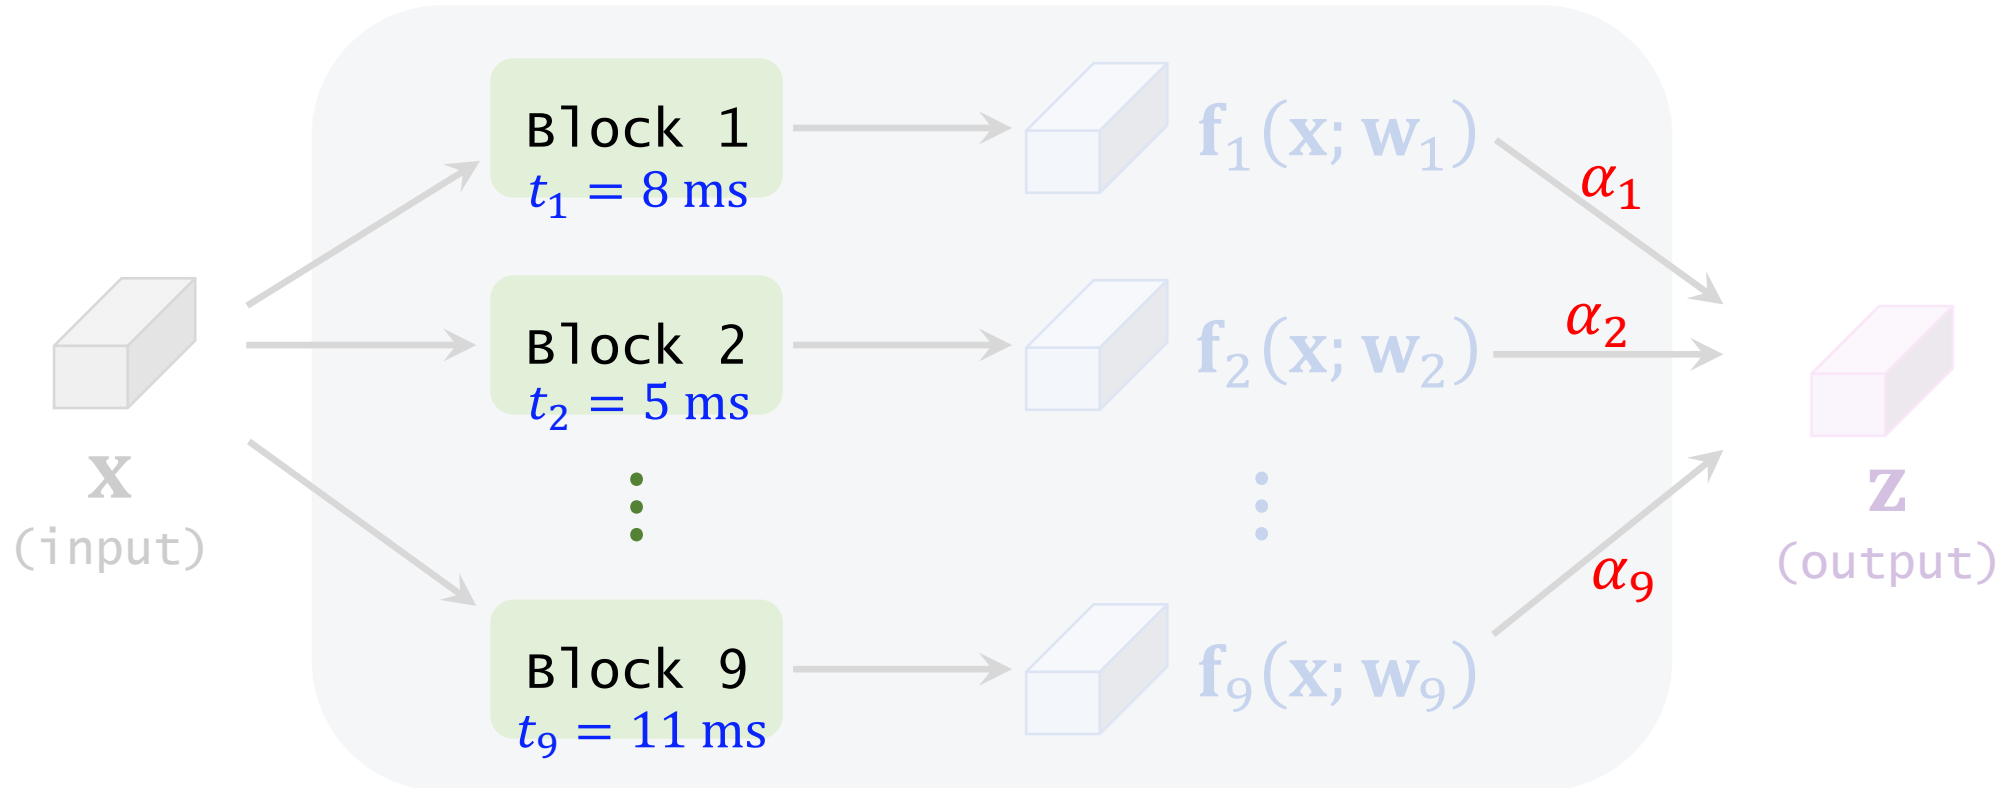
Latency

- Suppose the selected CNN will be deployed to iPhone 12.
- On iPhone 12, measure the latency caused by each block.



Latency

Weighted average of latencies: $\sum_{j=1}^9 t_j \cdot \alpha_j$.



Latency

- For layers $\underline{l = 1, \dots, 20}$ and blocks $\underline{j = 1, \dots, 9}$:

- Denote the measured latency (ms) by $t_j^{(l)}$.

- Denote the weights by $\underline{\alpha_j^{(l)}} = \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$.

Latency

- For layers $l = 1, \dots, 20$ and blocks $j = 1, \dots, 9$:

- Denote the measured latency (ms) by $t_j^{(l)}$.

- Denote the weights by $\alpha_j^{(l)} = \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$.

- Define: $\text{Lat}(\Theta) = \sum_{l=1}^{20} \sum_{j=1}^9 t_j^{(l)} \cdot \alpha_j^{(l)}$
 $= \sum_{l=1}^{20} \sum_{j=1}^9 t_j^{(l)} \cdot \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}$

Latency


Latency caused by the 20 layers:

$$\underline{\text{Lat}(\Theta)} = \sum_{l=1}^{20} \sum_{j=1}^9 t_j^{(l)} \cdot \frac{\exp(\theta_j^{(l)})}{\sum_{k=1}^9 \exp(\theta_k^{(l)})}.$$

- Encourage $\text{Lat}(\Theta)$ to be small.
- Apply (add or multiply) $\text{Lat}(\Theta)$ to the loss function.
- $\text{Lat}(\Theta)$ is a differential function of the parameters, Θ .

Trade off **accuracy** and **latency**

- Additive:

$$\min_{\mathcal{W}, \Theta} \underbrace{\frac{1}{n} \sum_{i=1}^n \text{Loss}(\mathbf{y}_i, \mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta))}_{\text{accuracy}} + \lambda \cdot \underbrace{\text{Lat}(\Theta)}_{\text{latency}}.$$


Trade off **accuracy** and **latency**

- Additive:

$$\min_{\mathcal{W}, \Theta} \frac{1}{n} \sum_{i=1}^n \text{Loss}(\mathbf{y}_i, \mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)) + \lambda \cdot \text{Lat}(\Theta).$$

- Multiplicative [1]:

$$\min_{\mathcal{W}, \Theta} \frac{1}{n} \sum_{i=1}^n \text{Loss}(\mathbf{y}_i, \mathbf{p}(\mathbf{x}_i; \mathcal{W}, \Theta)) \cdot \log_{\lambda} [\text{Lat}(\Theta)].$$

Reference:

1. Wu et al. [FBNet: Hardware-aware efficient convnet design via differentiable neural architecture search](#). In *CVPR*, 2019.

Hardware Awareness

- Some candidate blocks are suitable for GPU, while some are suitable for mobile devices.
 - For example, a candidate block is too small to fit in GPU.
 - But it can make full use of the A14 processor on iPhone 12.
- The optimal architectures for GPU and iPhone 12 are different.
 - For a GPU, the optimal architecture contains **big Conv** layers (good for accuracy, bad for latency.)
 - For iPhone 12, the optimal architecture contains **DepthWise Conv** layers (bad for accuracy, good for latency.)

Summary

Differentiable Architecture Search

- DARTS [1] automatically search neural architectures.
- This lecture explains DARTS using the example of [2].
- The objective function is a differentiable function of the parameters, $\Theta = \{\theta_j^{(l)}\}$, that determine network architecture.

Reference:

1. Liu, Simonyan, & Yang. DARTS: Differentiable Architecture Search. In *ICLR*, 2019.
2. Wu et al. FBNet: Hardware-aware efficient convnet design via differentiable neural architecture search. In *CVPR*, 2019.

Candidate Blocks & Super-net

- User manually prepare some (e.g., 9) candidate blocks.
- User manually specify the number of layers (e.g., 20.)
- Build a **super-net**: 20 layers; each layer contains the 9 parallel blocks.
- The output of each layer is the weighted sum of the 9 blocks; the weights are $\{\alpha_j^{(l)}\}$.

Candidate Blocks & Super-net


- Train the super-net (on the training set) to find the weights, $\alpha_j^{(l)}$ (for blocks $j = 1, \dots, 9$ and layers $l = 1, \dots, 20$).
- For the l -th layer, select the one among the 9 candidate blocks that has the biggest weight:

$$\operatorname{argmax}_{j \in \{1, \dots, 9\}} \alpha_j^{(l)} .$$


- The selected architecture has 20 layers, and each layer is one of the 9 candidate blocks.

Graph Perspective

Layer 1




Layer 20




Graph Perspective

Layer 1



Layer 20



Take efficiency into account

- Measure the **latency** (i.e., runtime of prediction making) caused by each of the $9 \times 20 = 180$ blocks.
- Take the weighted average (weights: $\alpha_j^{(l)}$) of the measured latencies for the 9 blocks in the l -th layer.
- $\text{Lat}(\Theta)$: sum of the **latencies** across the 20 layers.
- Apply (add or multiply) $\text{Lat}(\Theta)$ to the loss function.

Thank You!

<http://wangshusen.github.io/>