

量化计算怎么写？

5.1.1 整数运算

Integral Computing

```
float c = a * b;
```

```
int    c = a * b;
```

```
char   c = a * b;
```

5.1.1 整数运算

Integral Computing



float c = a * b;

int c = a * b;

char c = a * b;

指令	微指令个数	指令延迟	指令吞吐量
ADD	1	1	4
ADDSS	1	3	2
PADDB	1	1	4
MUL(r8)	1	3	1
MUL(r32)	2	3	1
MULPS	1	3	2
DIVSS	1	10.5	0.3

5.1.2 运算与访存

Computing & Memory Access



`float c = a * b;` 读a, b 写c, 总计 96 bit

`int c = a * b;` 读a, b 写c, 总计 96 bit

`char c = a * b;` 读a, b 写c, 总计 24 bit

5.1.2 运算与访存

Computing & Memory Access



`float c = a * b;` 读a, b 写c, 总计 96 bit

`int c = a * b;` 读a, b 写c, 总计 48 bit

`char c = a * b;` 读a, b 写c, 总计 24 bit

操作	时间代价
取指	100 (访主存) 1 (访缓存)
译码	>1
访存	100 (访主存) 1 (访缓存)
执行	1~5
写回	100 (访主存) 1 (访缓存)

5.1.3 向量化运算

SIMD Computing



在处理器中，为了加速大规模运算，通常会设计专用的向量化运算指令或向量化处理器。

典型地，在CPU中提供AVX512指令集，可以一次性处理512bit的数据。GPU中提供TensorCore，也可以一次性处理大量数据。

5.1.3 向量化运算

SIMD Computing



在处理器中，为了加速大规模运算，通常会设计专用的向量化运算指令或向量化处理器。

典型地，在CPU中提供AVX512指令集，可以一次性处理512bit的数据。GPU中提供TensorCore，也可以一次性处理大量数据。



优夜291

酷睿i5



avx512只有一个用处：发热

[△ 举报](#) 24楼 2021-08-22 14:58 [回复](#)

5.1.3 向量化运算

SIMD Computing



```
char function(char* array) {  
    return (  
        array[0] + array[1] + array[2] + array[3] +  
        array[4] + array[5] + array[6] + array[7] +  
        array[8] + array[9] + array[10] + array[11] +  
        array[12] + array[13] + array[14] + array[15]  
    );  
}
```


5.1.3 向量化运算

SIMD Computing



```
char function(char* array) {  
    return (  
        array[0] + array[1] + array[2] + array[3] +  
        array[4] + array[5] + array[6] + array[7] +  
        array[8] + array[9] + array[10] + array[11] +  
        array[12] + array[13] + array[14] + array[15]  
    );  
}
```

6.1.1 量化乘法

Quantized Mul



```
void Mul(  
    float** input_a, float** input_b,  
    float** output, const unsigned int num_of_elements) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = input_a[i][j] * input_b[i][j];  
}
```

6.1.1 量化乘法

Quantized Mul



```
void Mul(  
    char** input_a, char** input_b,  
    char** output, const unsigned int num_of_elements,  
    const float scale_a, const float scale_b, const float scale_c) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = input_a[i][j] * input_b[i][j];  
}
```

6.1.1 量化乘法

Quantized Mul

4.1 量化算子

Quantize Function

```
float value = 1.0; float scale = 0.1;
```

```
int qt32 = round_fn(value / scale);
```

```
char qt8 = clip(qt32, Q_MIN, Q_MAX)
```

```
for (unsigned int j = 0; j < num_of_elements; j++)
```

```
output[i][j] = input_a[i][j] * input_b[i][j];
```

```
}
```

4.1.3 反量化算子

Dequantize Function

```
char value = 1; float scale = 0.1;
```

```
float deq = (value * scale);
```

6.1.1 量化乘法

Quantized Mul



```
void Mul(  
    char** input_a, char** input_b,  
    char** output, const unsigned int num_of_elements,  
    const float scale_a, const float scale_b, const float scale_c) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = input_a[i][j] * scale_a * input_b[i][j] * scale_b / scale_c;  
}
```

6.1.1 量化乘法

Quantized Mul



```
void Mul(  
    char** input_a, char** input_b,  
    char** output, const unsigned int num_of_elements,  
    const float scale_a, const float scale_b, const float scale_c) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = clip(round_fn(input_a[i][j] * input_b[i][j] *  $\frac{s_a s_b}{s_c}$ ));  
}
```

6.1.1 量化乘法

Quantized Mul



```
void Mul(  
    char** input_a, char** input_b,  
    char** output, const unsigned int num_of_elements,  
    const float scale_a, const float scale_b, const float scale_c) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = clip(round_fn((input_a[i][j] * input_b[i][j]) << round(log2  $\frac{s_a s_b}{s_c}$ ))));  
}
```

6.1.1 量化乘法

Quantized Mul



```
output[i][j] =  
    (input_a[i][j] - offset_a) * scale_a *  
    (input_b[i][j] - offset_b) * scale_b / scale_c + offset_c;
```

在对称量化基础上，再引入 offset_a, offset_b, offset_c

6.1.1 量化乘法

Quantized Mul



output[i][j] =

(input_a[i][j] - offset_a) * scale_a *

(input_b[i][j] - offset_b) * scale_b / scale_c + offset_c;

记作:

$$c = (((b - o_b)(a - o_a) \frac{s_a s_b}{s_c})) + o_c$$

6.1.1 量化乘法

Quantized Mul



$$c = (((b - o_b)(a - o_a) \frac{s_a s_b}{s_c})) + o_c$$

展开得：

$$c = ab \frac{s_a s_b}{s_c} - a o_b \frac{s_a s_b}{s_c} - b o_a \frac{s_a s_b}{s_c} + o_a o_b \frac{s_a s_b}{s_c} + o_c$$

6.1.1 量化乘法

Quantized Mul



$$c = (((b - o_b)(a - o_a) \frac{s_a s_b}{s_c})) + o_c$$

展开得：

$$c = ab \frac{s_a s_b}{s_c} - a o_b \frac{s_a s_b}{s_c} - b o_a \frac{s_a s_b}{s_c} + o_a o_b \frac{s_a s_b}{s_c} + o_c$$

记作：

$$c = \text{rescale}((b - o_b)(a - o_a), s_a, s_b, s_c, o_c)$$

6.1.2 量化加法

Quantized Add



```
void Add(  
    char** input_a, char** input_b,  
    char** output, const unsigned int num_of_elements,  
    const float scale_a, const float scale_b, const float scale_c) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = (input_a[i][j] * scale_a + input_b[i][j] * scale_b) / scale_c;  
}
```

6.1.2 量化加法

Quantized Add



```
output[i][j] = (input_a[i][j] * scale_a + input_b[i][j] * scale_b) / scale_c;
```

```
output[i][j] = (input_a[i][j] *  $\frac{s_a}{s_c}$  + input_b[i][j] *  $\frac{s_b}{s_c}$ );
```

6.1.2 量化加法

Quantized Add



```
output[i][j] = (input_a[i][j] * scale_a + input_b[i][j] * scale_b) / scale_c;
```

要求 scale_a, scale_b 一致!

```
output[i][j] = (input_a[i][j] + input_b[i][j] *  $\frac{s_i}{s_c}$ );
```

```
c = rescale(a + b + o_b + o_a, s_i, 1, s_c, o_c)
```

6.1.3 量化激活函数

Quantized Activation



```
void Clip(  
    float** input, float** output, float min, float max,  
    const unsigned int num_of_elements) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
        {  
            output[i][j] = MAX(input[i][j], min);  
            output[i][j] = MIN(input[i][j], max);  
        }  
}
```

6.1.3 量化激活函数

Quantized Activation



```
void Clip(  
    char** input, char** output, float min, float max  
    const float in_scale, const float out_scale, const unsigned int num_of_elements) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
        {  
            output[i][j] = MAX(input[i][j] * in_scale, min) / out_scale;  
            output[i][j] = MIN(input[i][j] * in_scale, max) / out_scale;  
        }  
    }
```


6.1.3 量化激活函数

Quantized Activation



```
output[i][j] = MAX(input[i][j] * in_scale, min) / out_scale;
```

要求 in_scale, out_scale 一致!

```
output[i][j] = MAX(input[i][j] * scale, min) / scale;
```

6.1.3 量化激活函数

Quantized Activation



```
output[i][j] = MAX(input[i][j] * in_scale, min) / out_scale;
```

要求 in_scale, out_scale 一致!

```
output[i][j] = MAX(input[i][j] * scale, min) / scale;
```

```
output[i][j] = MAX(input[i][j], min / scale);
```

6.1.3 被动量化算子

Passive Quantized Operator



```
output[i][j] = MAX(input[i][j], min / scale);
```

式中 scale 与 min 参数无关，此时称 min 参数被动量化
常见的被动量化参数包括：

Bias(Gemm, Conv, Lstm),

min(Clip), max(Clip),

padding value(Pad),

6.1.3 被动量化算子

Passive Quantized Operator



```
output[i][j] = MAX(input[i][j], min / scale);
```

式中 scale 被输入和输出同时共享，此时算子的运算不改变量化参数，我们称这类算子为被动量化算子。常见的被动量化算子包括：

Pad, Clip, Relu, MaxPooling, Reshape, Concat, Split, Transpose, Slice, Permute

6.1.4 量化矩阵乘

Quantized Gemm



```
void MatMul(
    ELEMENT_TYPE** input, ELEMENT_TYPE** weight, ELEMENT_TYPE* bias,
    ELEMENT_TYPE** output, const unsigned int num_of_elements) {
```

```
    ACCUMULATOR_TYPE Accumulator[16];
```

```
    for (unsigned int i = 0; i < num_of_elements; i += 4) {
```

```
        // Pack A[i: i + 4][.], Send Packed A to L2
```

```
        ELEMENT_TYPE* packedA = LhsPackElement(input, num_of_elements, i);
```

```
        for (unsigned int j = 0; j < num_of_elements; j += 4) {
```

```
            // Pack B[.][j: j + 4], Send Packed B to L2
```

```
            ELEMENT_TYPE* packedB = RhsPackElement(weight, num_of_elements, j);
```

```
                for (unsigned int k = 0; k < num_of_elements; k += 1) {
```

```
                    // Accumulator = A[i: i + 4][.] * B[.][j: j + 4]
```

```
                    MatMul4x4(packedA, packedB, Accumulator, k);
```

```
                }
```

```
            for (unsigned int k = 0; k < 4; k += 1) {
```

```
                output[i + k][j + 0] = Accumulator[0] + bias[i + k];
```

```
                output[i + k][j + 1] = Accumulator[0] + bias[i + k];
```

```
                output[i + k][j + 2] = Accumulator[0] + bias[i + k];
```

```
                output[i + k][j + 3] = Accumulator[0] + bias[i + k];
```

```
            }
```

```
        }
```

```
    }
```

```
}
```

数据送上L2

改变数据排布

8bit向量化加速

8bit可以容纳更多数据

分块矩阵乘

8bit 向量化加速

结果为16 - 32位

Bias Add

32bit 向量化加速

结果为32 - 64位

Rescale 到 int8

6.1.4 量化矩阵乘

Quantized Gemm

```
// Pack A[i: i + 4][.], Send Packed A to L2
```

```
ELEMENT_TYPE* packedA = LhsPackElement(input, num_of_elements, i);
```

```
__declspec(noinline) ELEMENT_TYPE* LhsPackElement(  
    ELEMENT_TYPE** input, unsigned int num_of_element, unsigned int row) {  
    ELEMENT_TYPE* packed = new ELEMENT_TYPE[num_of_element * 4];  
    unsigned int k = 0;  
    for (unsigned int i = 0; i < num_of_element; i++) {  
        for (unsigned int j = 0; j < 4; j++) {  
            packed[k++] = input[i][row + j];  
        }  
    }  
    return packed;  
}
```

6.1.4 量化矩阵乘

Quantized Gemm

```
// Accumulator = A[i: i + 4][.] * B[.][j: j + 4]
```

```
MatMul4x4(packedA, packedB, Accumulator, k);
```

```
__declspec(noinline) void MatMul4x4(  
    ELEMENT_TYPE* packedA, ELEMENT_TYPE* packedB,  
    ACCUMULATOR_TYPE* accumulator, unsigned int offset) {  
    // 所有计算仅操作下列元素  
    // PackedA[offset: offset + 16]  
    // PackedB[offset: offset + 16]  
    // C[16], 共计48个, 可以全部送入寄存器与L1 Cache  
    accumulator[0] = packedA[0] * packedB[0] + accumulator[0];  
    accumulator[1] = packedA[0] * packedB[1] + accumulator[1];  
    accumulator[2] = packedA[0] * packedB[2] + accumulator[2];  
    accumulator[3] = packedA[0] * packedB[3] + accumulator[3];  
    // ...  
}
```


6.1.4 量化矩阵乘

Quantized Gemm

```
output[i + k][j + 0] = Accumulator[0] + bias[i + k];
```

```
accumulator[0] = packedA[0] * packedB[0] + accumulator[0];
```

```
output[i + k][j + 0] = packedA[0] * scale_A * packedB[0] * scale_B + bias[i + k];
```

要求 `bias_scale` 与 `scale_A * scale_B` 一致!

猜猜看 `bias_offset` 是谁?

6.1.4 量化矩阵乘

Quantized Gemm

```
output[i + k][j + 0] = Clip(round(packedA[0] * packedB[0] + bias[i + k]) *  $\frac{s_a s_b}{s_c}$ ));
```

1. int8 矩阵乘在累加器中运行，结果为int32
2. 处理bias_add，结果为int32
3. 执行rescale，结果为fp32
4. 取整，截断，结果为int8

6.1.5 量化非线性运算

Quantized Non - Linear Function



算子诸如：Exp, Tanh, Sigmoid, Softmax, Swish, Resize，内部包含非线性运算。

不可以直接量化，在不同处理器上做法不同：

在CPU, GPU上，这类算子的计算不量化，以全精度模式运行。

在FPGA, ASIC, DSP上，需要更改算子计算逻辑，以线性运算拟合或直接查表。

6.1.5 量化非线性运算

Quantized Non - Linear Function



$$\exp(x) = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + R_4$$

1. 以泰勒展开方式进行拟合
2. 对结果进行rescale，从而产生 int8 的结果

请注意，exp的数值范围极广，不利于量化

6.1.5 量化非线性运算

Quantized Non - Linear Function



$$\exp(x, \text{shift}) = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + R_4$$



4.1.4 整数量化

Power - of - 2 Quantization

```
float value = 1.0; int shift = 1;  
int qt32 = round_fn(value * (2 << shift)) ;  
char qt8 = clip(qt32, Q_MIN, Q_MAX)
```

6.1.5 量化非线性运算

Quantized Non - Linear Function



$$\exp(x, \text{shift}) = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + R_4$$

x	shift	output(int32)	output(fp32)
1	0	3	2.71828
2	0	7	2.71828^2
3	0	20	2.71828^3
4	0	54	2.71828^4

对结果进行rescale，从而产生 int8 的结果

6.1.5 量化非线性运算

Quantized Non - Linear Function



```
template<typename Dtype>
void cuda_sigmoid_table_lookup(const int N,
                               const Dtype* input,
                               const Dtype* table,
                               Dtype* output,
                               int fragpos);
```

```
template<typename Dtype>
void cuda_sigmoid_simulation(const int N,
                             const Dtype* input,
                             Dtype* output);
```

```
template<typename Dtype>
void cuda_tanh_simulation(const int N,
                          const Dtype* input,
                          Dtype* output);
```

```
template<typename Dtype>
void cuda_tanh_table_lookup(const int N,
                             const Dtype* input,
                             const Dtype* table,
                             Dtype* output,
                             int fragpos);
```

```
template<typename Dtype>
__global__ static void _sigmoid_simulation(const int N,
                                           const Dtype fuzz,
                                           const Dtype* input,
                                           Dtype* output) {
    NNDCT_KERNEL_LOOP(i, N){
        if (input[i] >= 8.0)
            output[i] = 1.0 - fuzz;
        else if (input[i] < -8.0)
            output[i] = 0.0;
        else {
            int x = int(input[i] * pow(2, 9));
            output[i] = sigmoid_short_sim(x, 9, 8) / (pow(2.0, 8));
        }
    }
}
```

口诀



量化计算量化算，中间结果精度高
中间算完转尺度，转完尺度取整数
加法减法不能转，被动算子也一样
非线性函数查表算，不然你就等死吧


6.1.7 量化复杂算子

Quantized Complex Operation



联系我们

<https://github.com/openppl-public>



广告位招租



微信群



QQ群 (入群密令OpenPPL)