

神经网络图优化与量化

6.2.1 啥是计算图



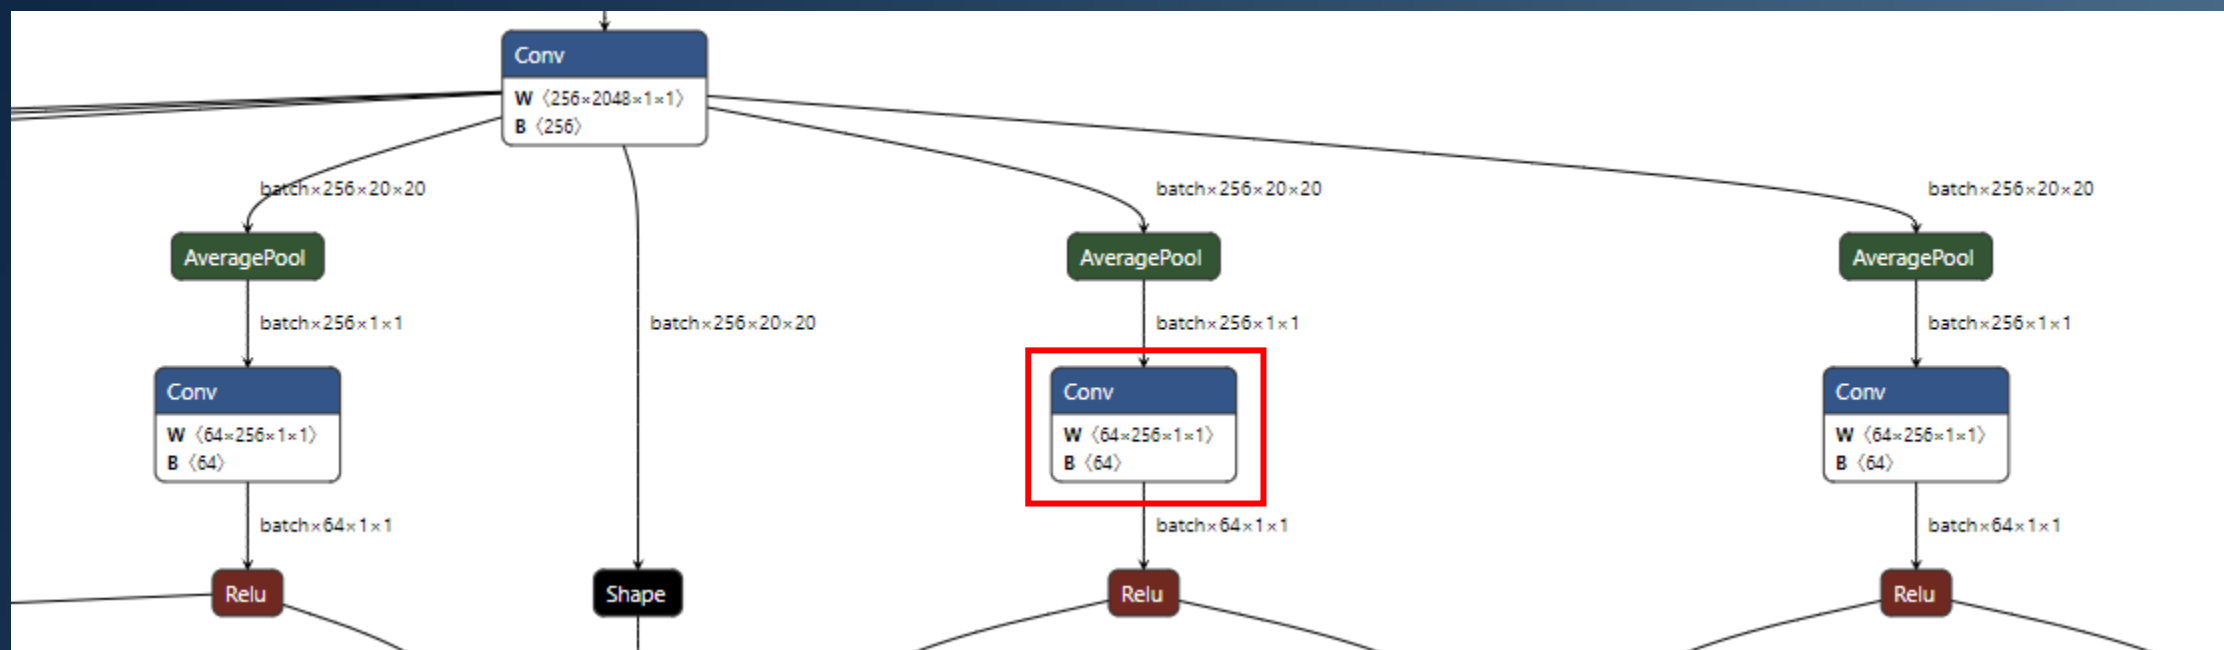
What is a Computational Graph.

$C = \{N, E, I, O\}$ 一个计算图可以表示为一个由节点、边集、输入边、输出边组成的四元组。

- 计算图是一个有向联通无环图，其中节点也被称作为算子。
- 算子必定有边相连，输入边，输出边不为空。
- 计算图中可以有重边。

6.2.1 算子

What is a Computational Graph.



6.2.1 算子

What is a Computational Graph.


Operator	Since version	
Abs	13, 6, 1	Constant 13, 12, 11, 9, 1
Acosh	9	ConstantOfShape 9
Add	14, 13, 7, 6, 1	Conv 11, 1
And	7, 1	ConvInteger 10
ArgMax	13, 12, 11, 1	ConvTranspose 11, 1
ArgMin	13, 12, 11, 1	Cos 7
Asin	7	Cosh 9
Asinh	9	CumSum 14, 11
Atan	7	DepthToSpace 13, 11, 1
Atanh	9	DequantizeLinear 13, 10
AveragePool	11, 10, 7, 1	Det 11
BatchNormalization	15, 14, 9, 7, 6, 1	Div 14, 13, 7, 6, 1
		Dropout 13, 12, 10, 7, 6, 1

<https://github.com/onnx/onnx/blob/main/docs/Operators.md>

6.2.1 算子

What is a Computational Graph.

算子是神经网络的最小调度单位，但很遗憾的是，它并不是原子的：一个复杂的算子可以被更细粒度的算子所表示：



我们总是以算子为单位去支持你的网络。

在你的网络中你应该尽量避免使用特殊算子。

6.2.2 算子融合加速

Graph Fusion.

```
__declspec(noinline) void MatMul(
    ELEMENT_TYPE** input, ELEMENT_TYPE** weight,
    ELEMENT_TYPE** output, const unsigned int num_of_elements) {
    for (unsigned int i = 0; i < num_of_elements; i++)
        for (unsigned int j = 0; j < num_of_elements; j++)
            for (unsigned int k = 0; k < num_of_elements; k++)
                output[i][j] += input[i][k] * weight[k][j];
}
```

```
__declspec(noinline) void BiasAdd(
    ELEMENT_TYPE** input, ELEMENT_TYPE* bias,
    ELEMENT_TYPE** output, const unsigned int num_of_elements) {
    for (unsigned int i = 0; i < num_of_elements; i++)
        for (unsigned int j = 0; j < num_of_elements; j++)
            output[i][j] += bias[i];
}
```

6.2.2 算子融合加速

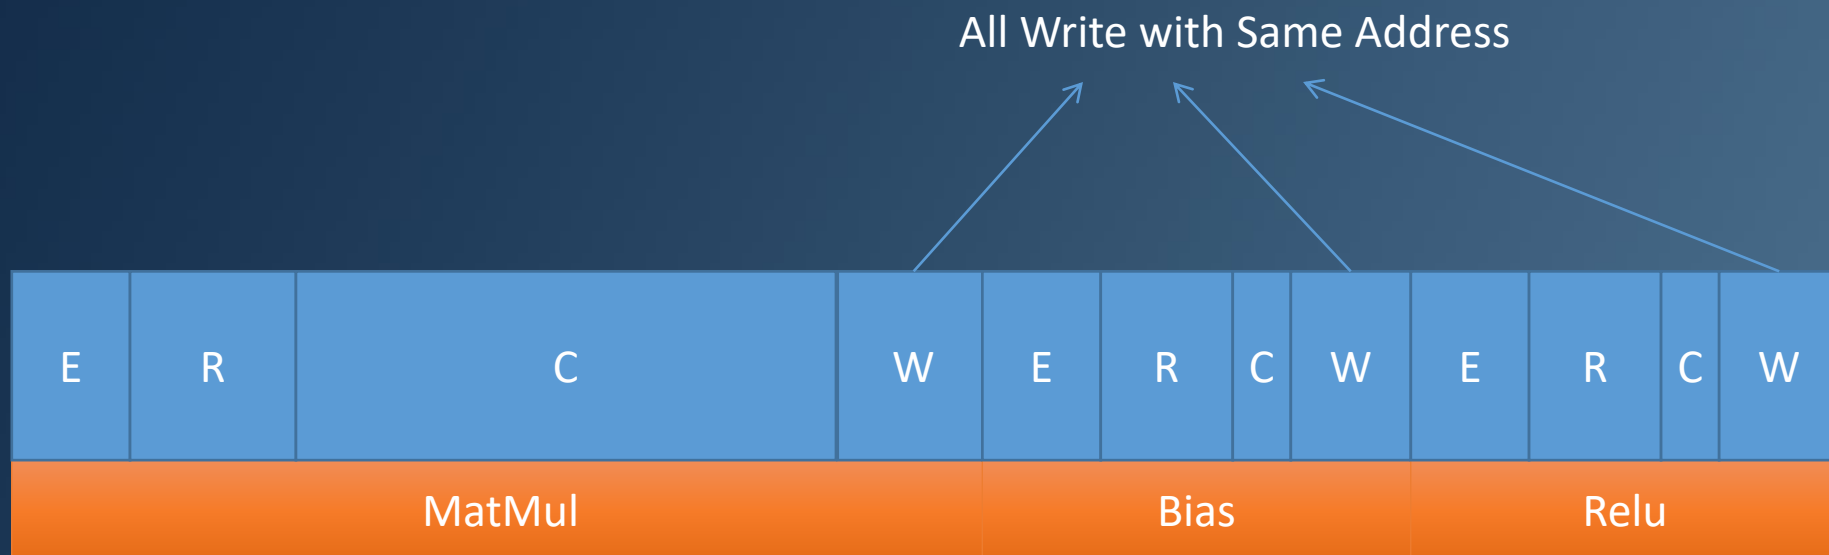
Graph Fusion.

```
__declspec(noinline) void Relu(  
    ELEMENT_TYPE** input, ELEMENT_TYPE** output,  
    const unsigned int num_of_elements) {  
    for (unsigned int i = 0; i < num_of_elements; i++)  
        for (unsigned int j = 0; j < num_of_elements; j++)  
            output[i][j] = input[i][j] * (input[i][j] > 0);  
}
```

在Matmul + Bias + Relu的子网中，如果不融合算子，output将至少被写入3次并且启动3个算子的速度也不是很快。

6.2.2 算子融合加速

Graph Fusion.



E: Task Emission

R: Read

C: Computing

W: Write

6.2.2 算子融合加速

Graph Fusion.

```
__declspec(noinline) void MatMul(
    ELEMENT_TYPE** input, ELEMENT_TYPE** weight, ELEMENT_TYPE* bias,
    ELEMENT_TYPE** output, const unsigned int num_of_elements) {
    for (unsigned int i = 0; i < num_of_elements; i++){
        for (unsigned int j = 0; j < num_of_elements; j++){
            int accumulator = 0;
            for (unsigned int k = 0; k < num_of_elements; k++){
                accumulator += input[i][k] * weight[k][j];
            }
            output[i][j] = accumulator + bias[j] > 0 ? accumulator + bias[j]: 0;
        }
    }
}
```

6.2.2 算子融合加速

Graph Fusion.



图融合减少访存，以及算子overhead

6.2.3 常见计算图优化

Widely - used Graph Optimization

6.2.3.1 激活函数融合

6.2.3.2 移除 Batchnorm 与 Dropout


6.2.3.3 常量折叠

6.2.3.4 矩阵乘融合

6.2.3.5 Conv - Add 融合

6.2.3 常见计算图优化

Widely - used Graph Optimization



常见计算算子 : Conv, ConvTranpose, Gemm

常见激活函数 : Relu, Clip(Relu6), PRelu, Tanh, Sigmoid, Swish

```
output[i][j] = activation_fn(accumulator + bias[j]);
```


6.2.3 常见计算图优化

Widely - used Graph Optimization



6.2.3 常见计算图优化

Widely - used Graph Optimization



计算算子 : $Y = WX + B$


Batchnorm : $Y' = \text{gamma} * \frac{Y - \text{mean}}{\text{var}} + \text{beta}$

Merged : $Y' = \text{gamma} * \frac{WX + B - \text{mean}}{\text{var}} + \text{beta} = \frac{\text{gamma}}{\text{var}} WX + \frac{\text{gamma}}{\text{var}} (B - \text{mean}) + \text{beta}$

6.2.3 常见计算图优化


Widely - used Graph Optimization

```
alpha = bn_op.parameters[0].value
beta  = bn_op.parameters[1].value
mean  = bn_op.parameters[2].value
var    = bn_op.parameters[3].value
if computing_op.type == 'Conv':
    # calculate new weight and bias
    scale = alpha / np.sqrt(var + epsilon)
    w = w * scale.reshape([-1, 1, 1, 1])
    b = alpha * (b - mean) / np.sqrt(var + epsilon) + beta
elif computing_op.type == 'Gemm':
    # calculate new weight and bias
    scale = alpha / np.sqrt(var + epsilon)
    if computing_op.attributes.get('transB', 0): w = w * scale.reshape([-1, 1])
    else: w = w * scale.reshape([1, -1])
    b = alpha * (b - mean) / np.sqrt(var + epsilon) + beta
```



6.2.3 常见计算图优化

Widely - used Graph Optimization




Add1 : $Y = X + 1$

Add2 : $Y' = Y + 1$

Merged : $Y' = X + 2$

6.2.3 常见计算图优化

Widely - used Graph Optimization



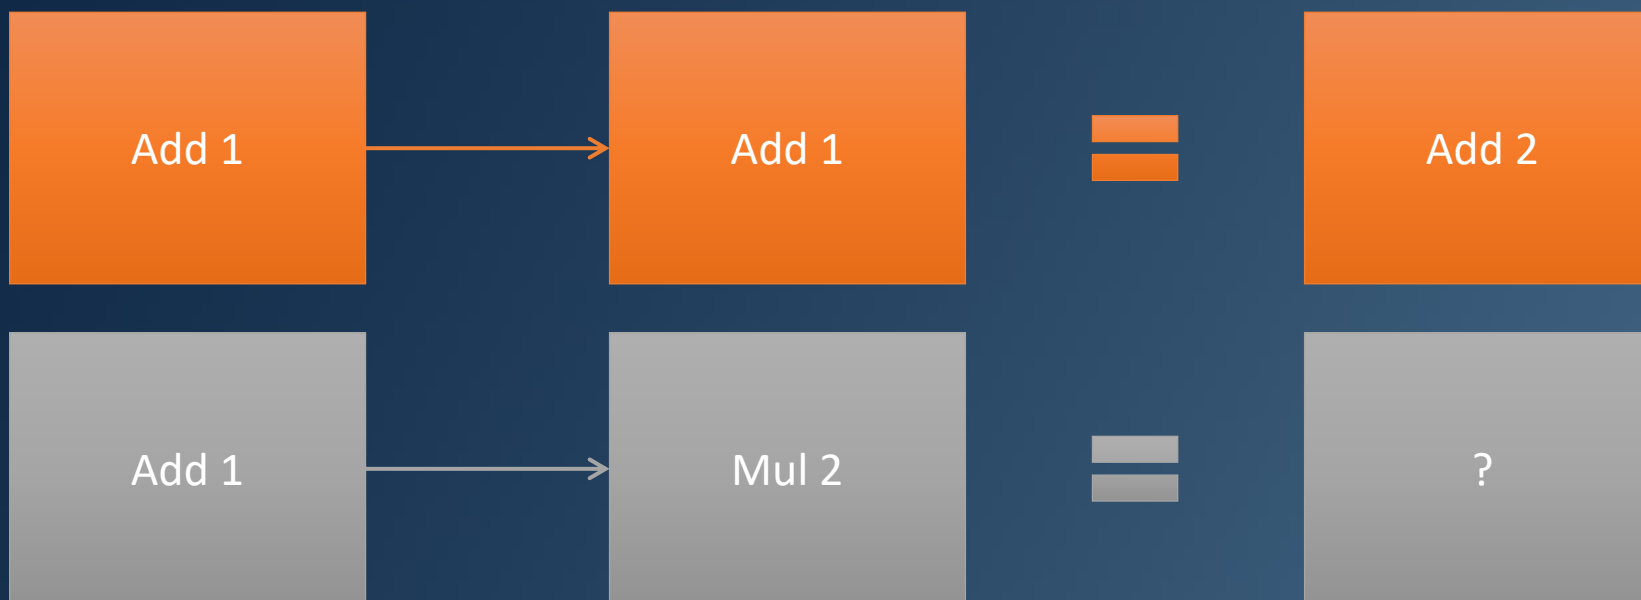
Add1 : $Y = X + 1$

Mul2 : $Y' = Y * 2$

Merged : $Y' = (X + 1) * 2 = X * 2 + 2$

6.2.3 常见计算图优化

Widely - used Graph Optimization



Add1 : $Y = X + 1$

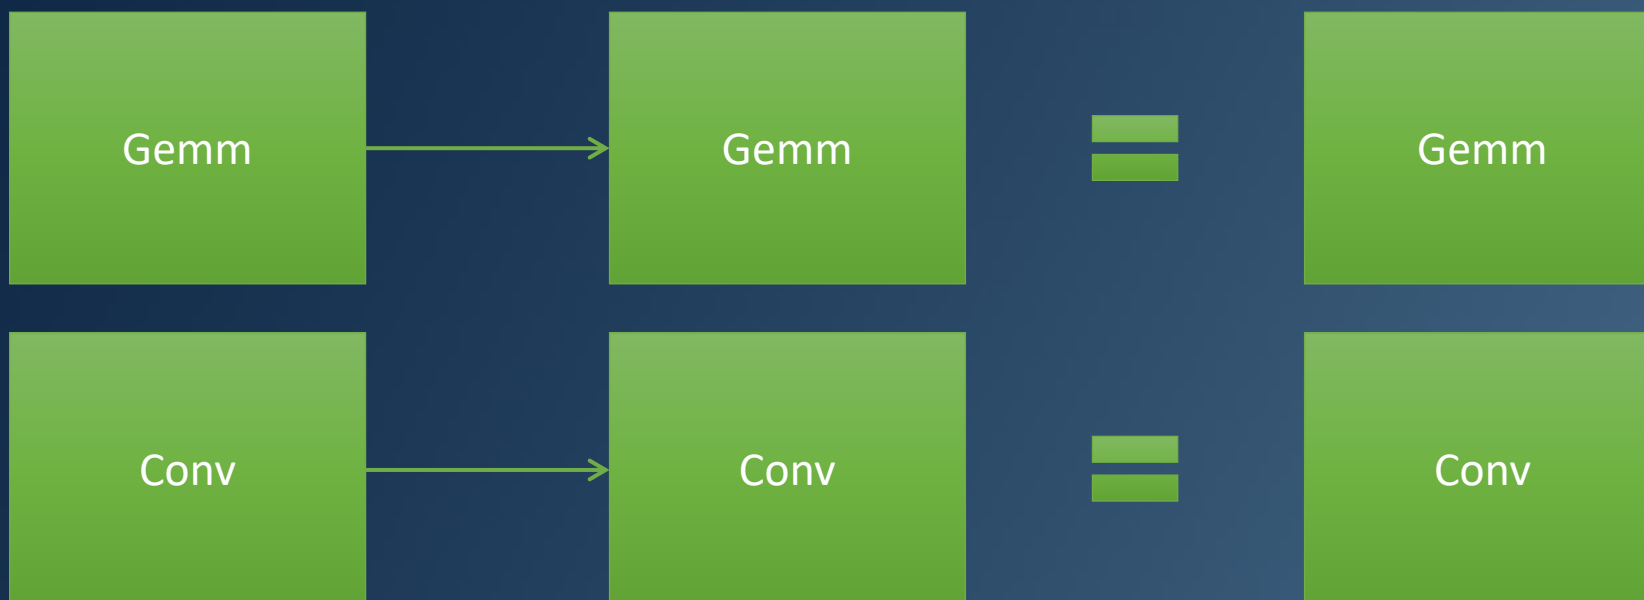
Mul2 : $Y' = Y * 2$

Merged : $Y' = (X + 1) * 2 = X * 2 + 2?$

EMM... WE DO NOT HAVE AN OP TO RUN $X*2+2$!

6.2.3 常见计算图优化

Widely - used Graph Optimization



计算算子1 : $Y = W_1X + B_1$

计算算子2 : $Y' = W_2Y + B_2$

融合后 : $Y' = W_2(W_1X + B_1) + B_2$

6.2.3 常见计算图优化



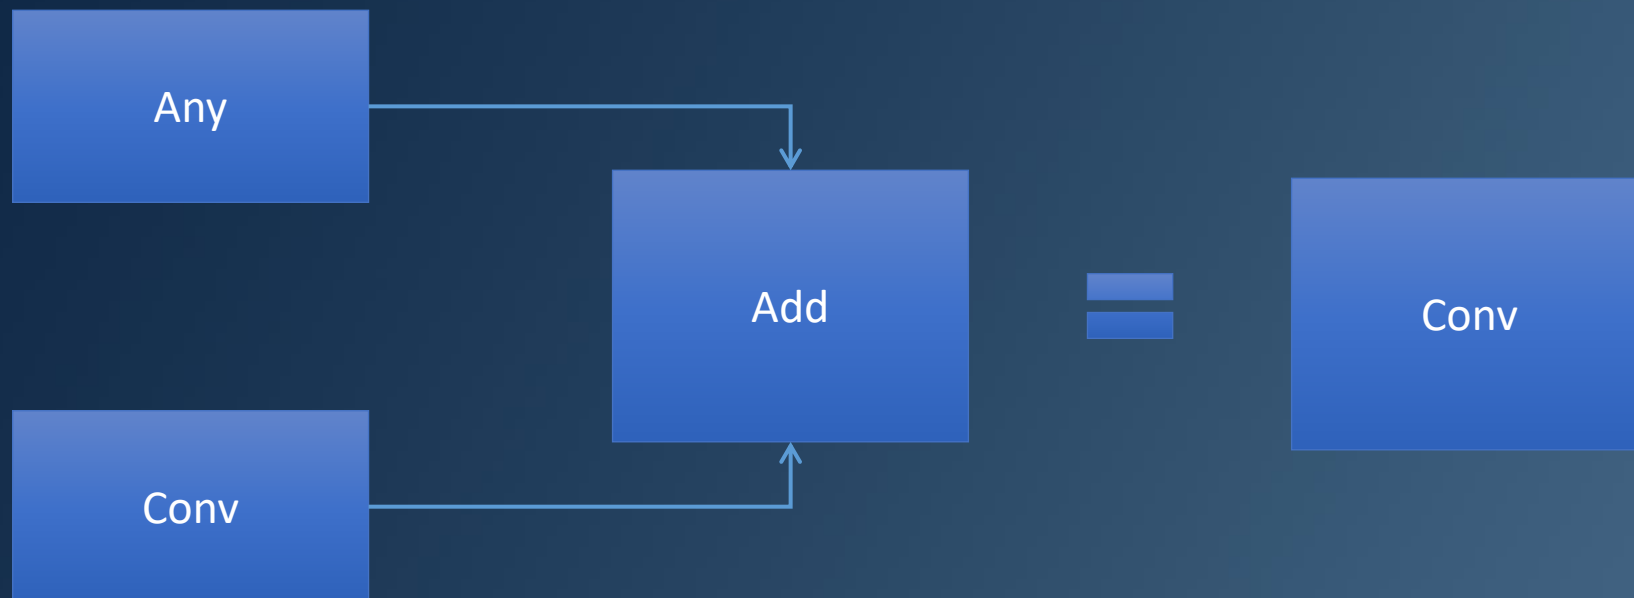
Widely - used Graph Optimization

```
def svd_for_factorization(self, w: torch.Tensor):  
    assert w.ndim == 2  
    u, s, v = torch.svd(w)  
    a = torch.matmul(u, torch.diag(torch.sqrt(s)))  
    b = torch.matmul(torch.diag(torch.sqrt(s)), v.transpose(0, 1))  
    print(a.max(), b.max(), w.max())  
    return a, b
```

```
if operation.type == 'Gemm':  
    w = operation.parameters[0].value  
    w = w.transpose(0, 1)  
    if self.method == 'svd':  
        a, b = self.svd_for_factorization(w)
```

6.2.3 常见计算图优化

Widely - used Graph Optimization



$$\text{Conv} : Y_1 = W_1 X_1 + B_1$$

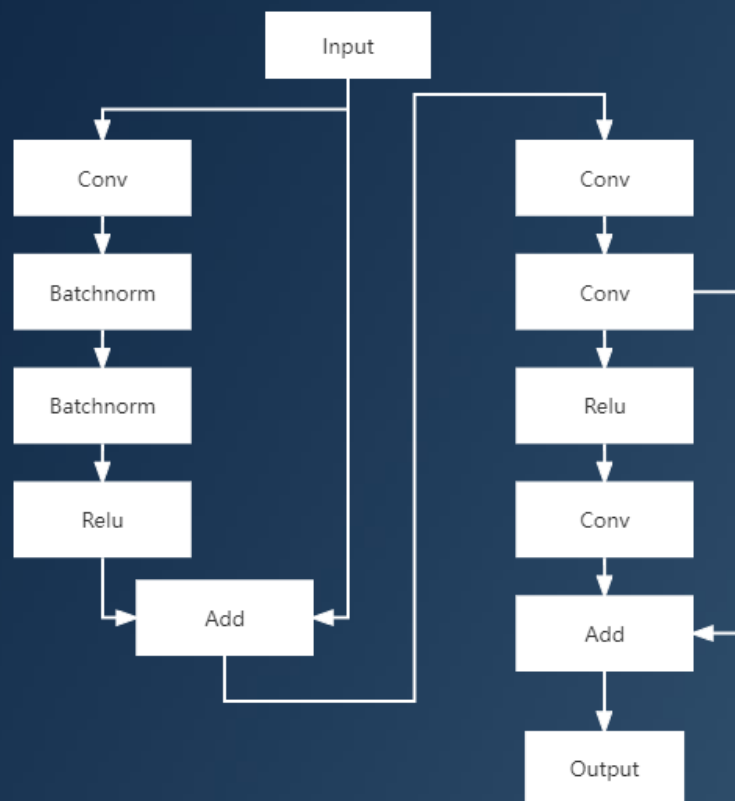
$$\text{Any} : Y_2$$

$$Y = Y_1 + Y_2$$

$$\text{融合后} : Y = W_1 X_1 + (Y_2 + B_1)$$

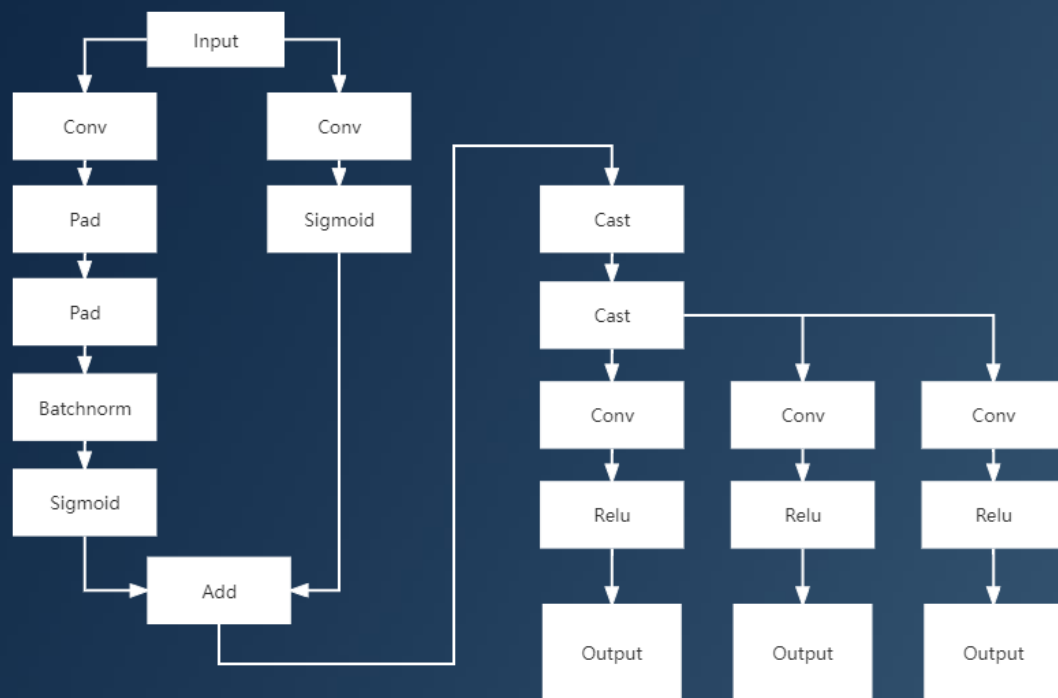
6.2.3 常见计算图优化

Widely - used Graph Optimization




6.2.3 常见计算图优化

Widely - used Graph Optimization



6.2.4 联合定点

Union - Quantize



```
class TensorQuantizationConfig(Serializable):  
    self._policy = policy  
    self._num_of_bits = num_of_bits  
    self._scale = scale  
    self._offset = offset  
    self.state = state  
    self._rounding = rounding  
    self._quant_min = quant_min  
    self._quant_max = quant_max  
    self._father_config = self # union - find
```

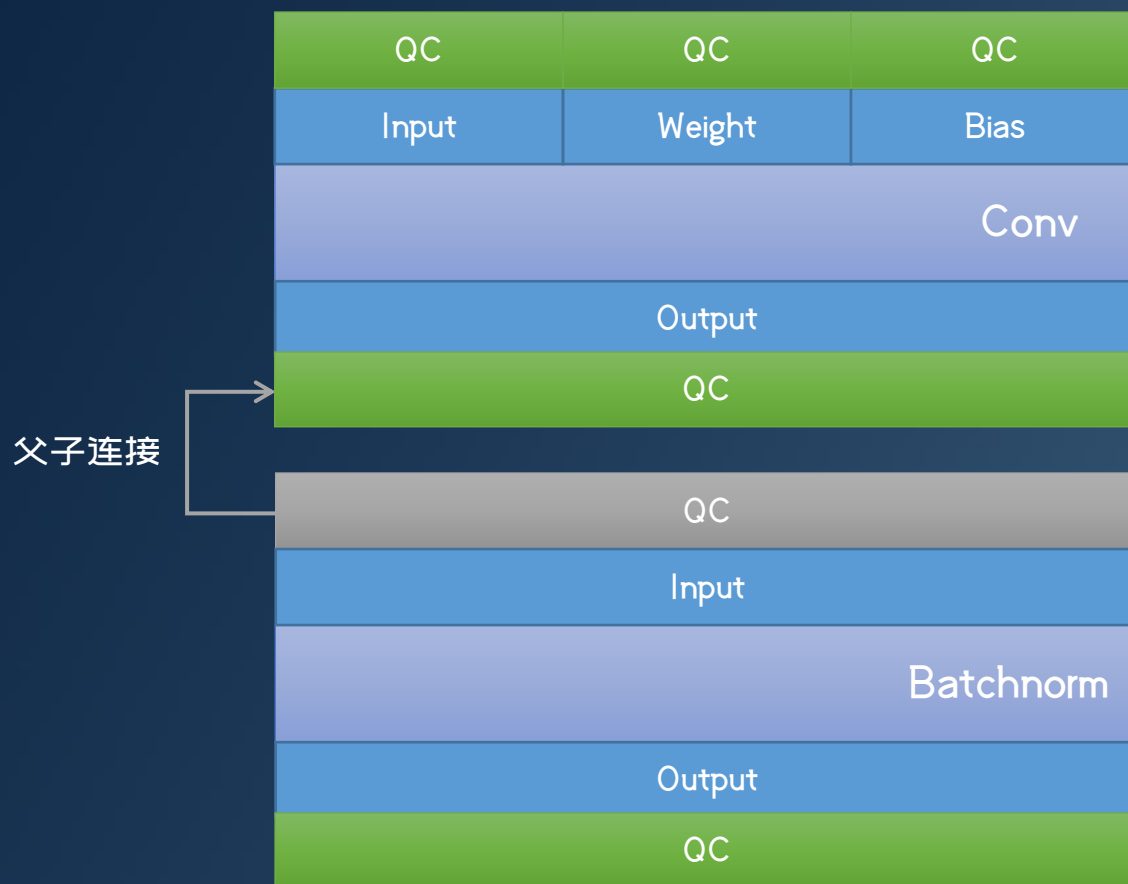

6.2.4 联合定点

Union - Quantize



6.2.4 联合定点

Union - Quantize



输入已经被量化，定点信息停用

6.2.4 联合定点

Union - Quantize

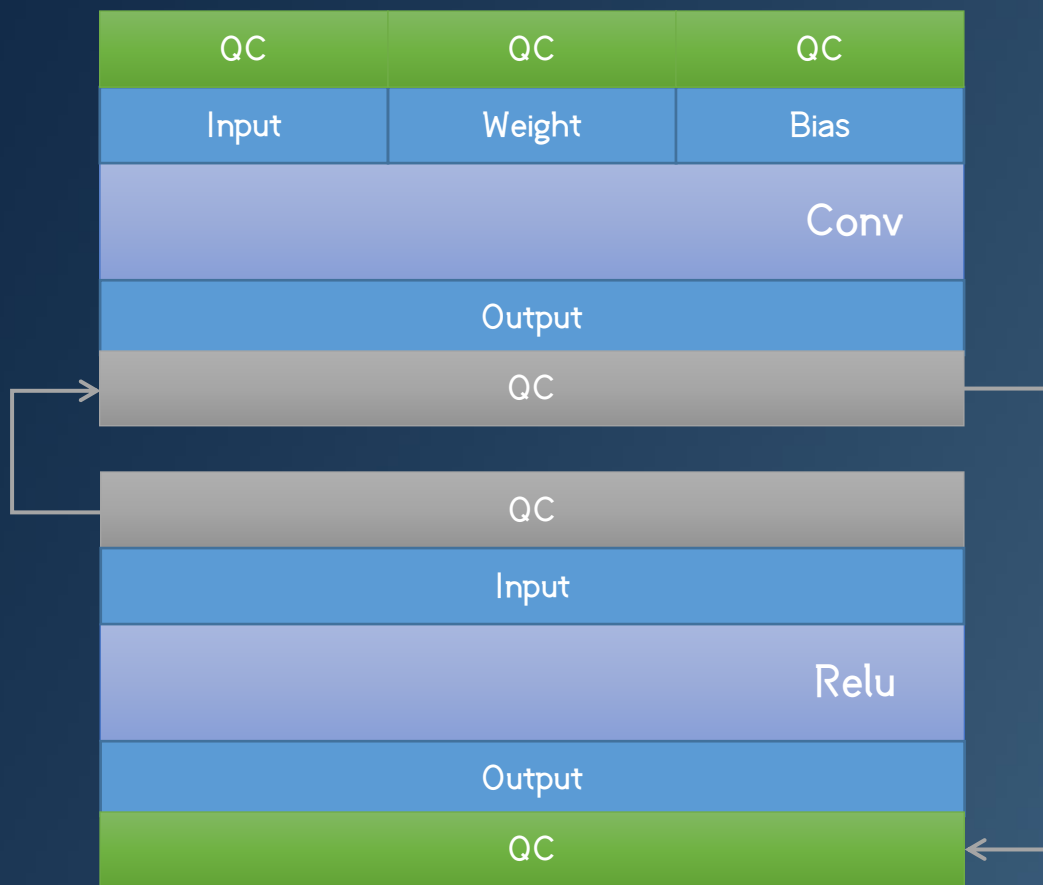


考虑到Conv与Batchnorm图融合
输出定点被停用

级联父子连接（并查集）

6.2.4 联合定点

Union - Quantize



类似地，激活函数联合定点

6.2.4 联合定点

Union - Quantize

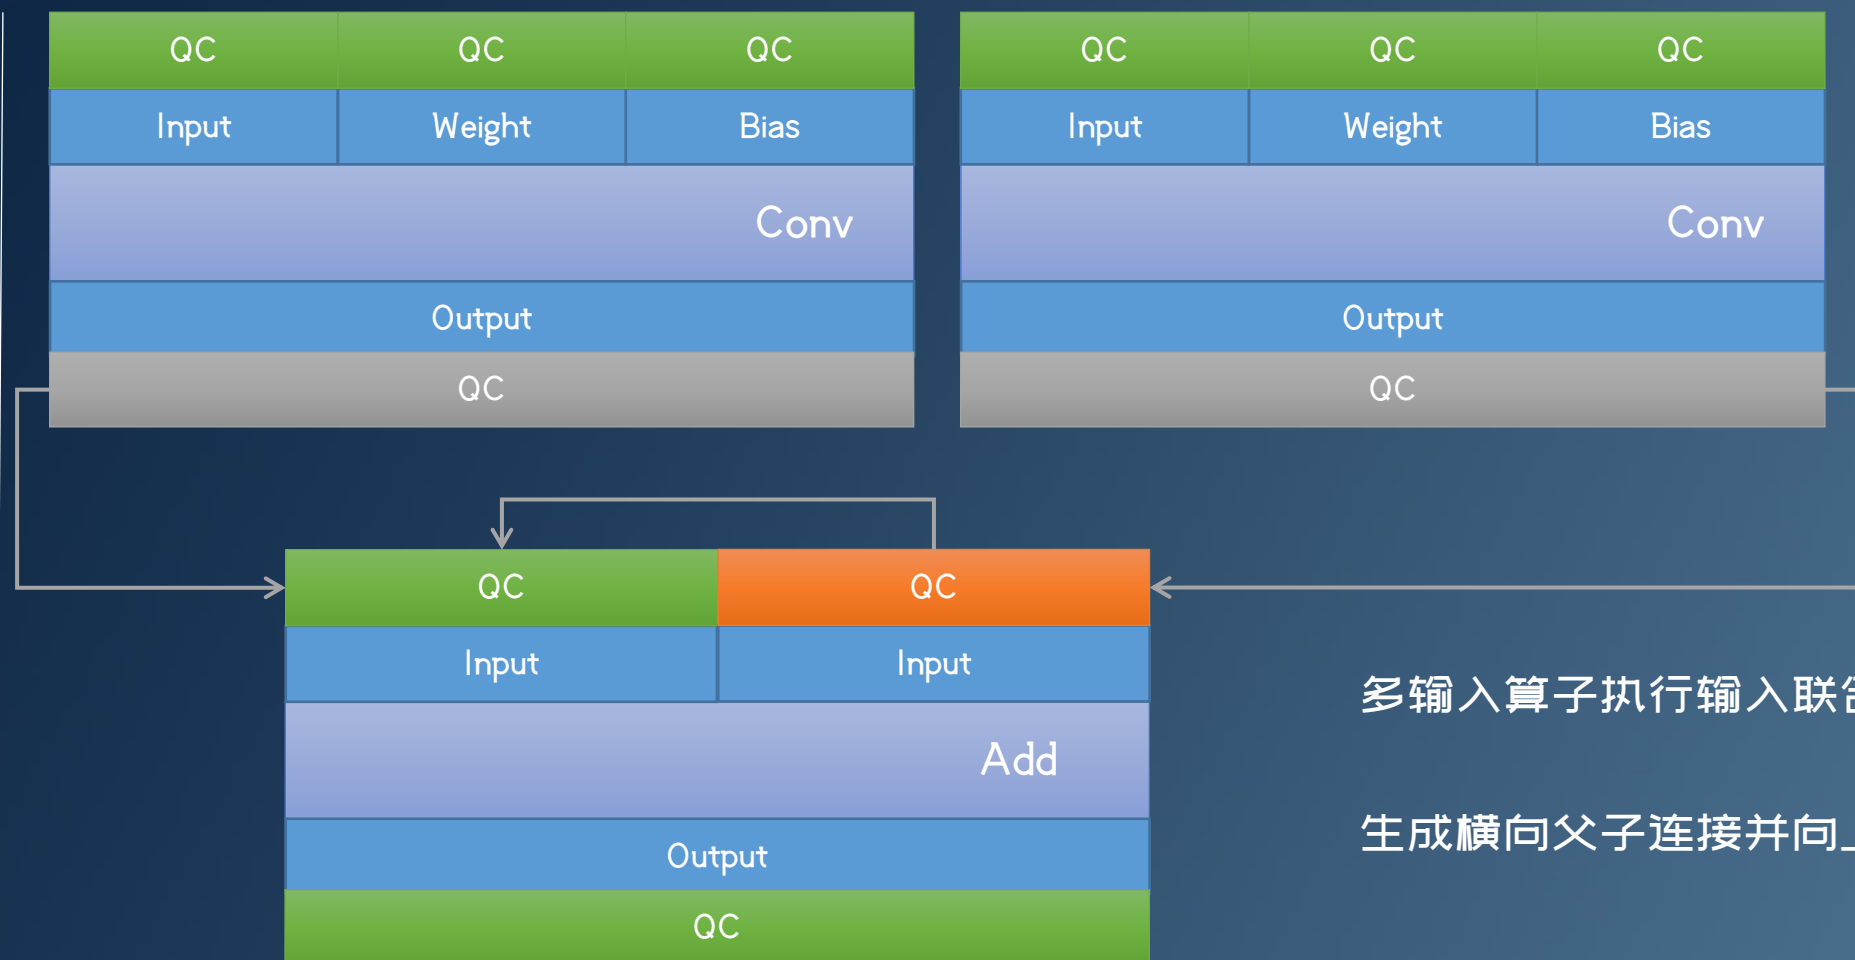


被动算子量化定点信息全部停用

向上生成级联父子连接

6.2.4 联合定点

Union - Quantize

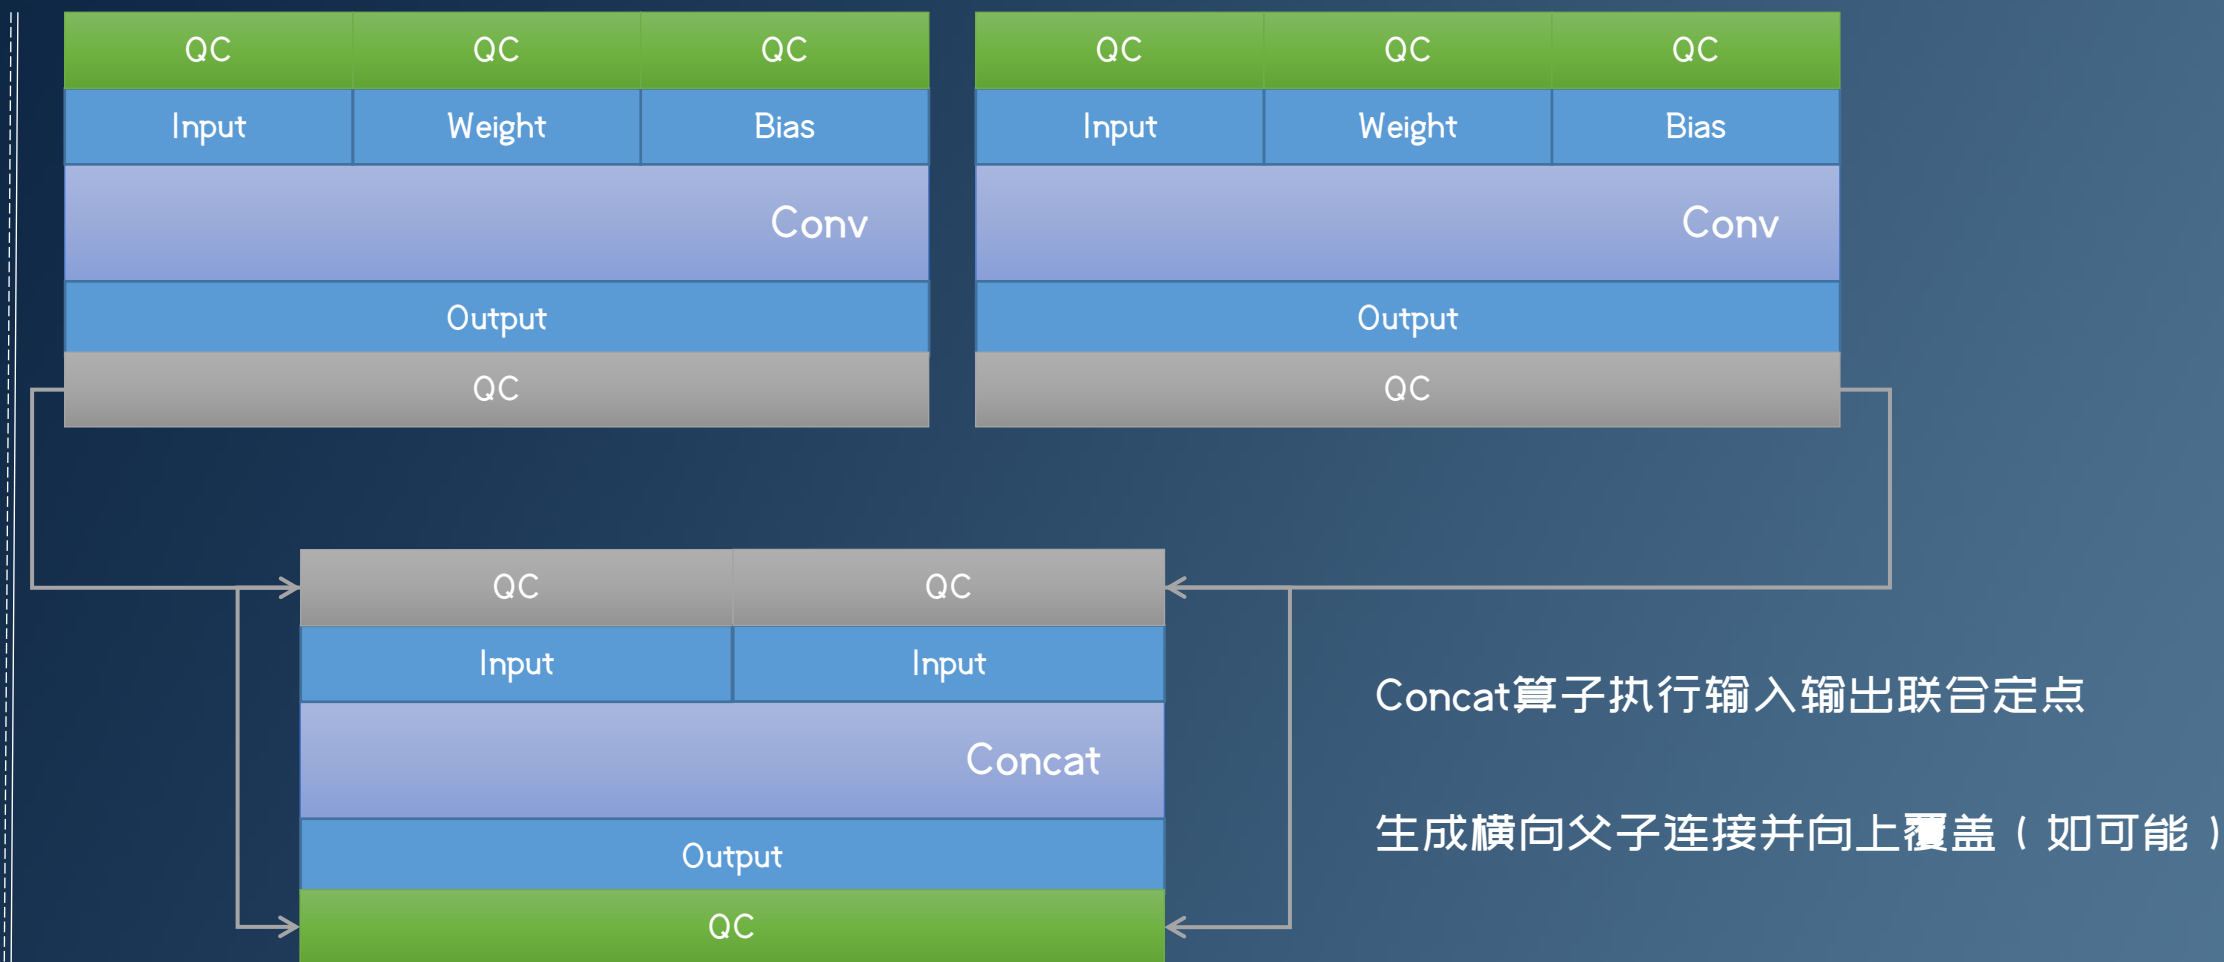


多输入算子执行输入联合定点

生成横向父子连接并向上覆盖（如可能）

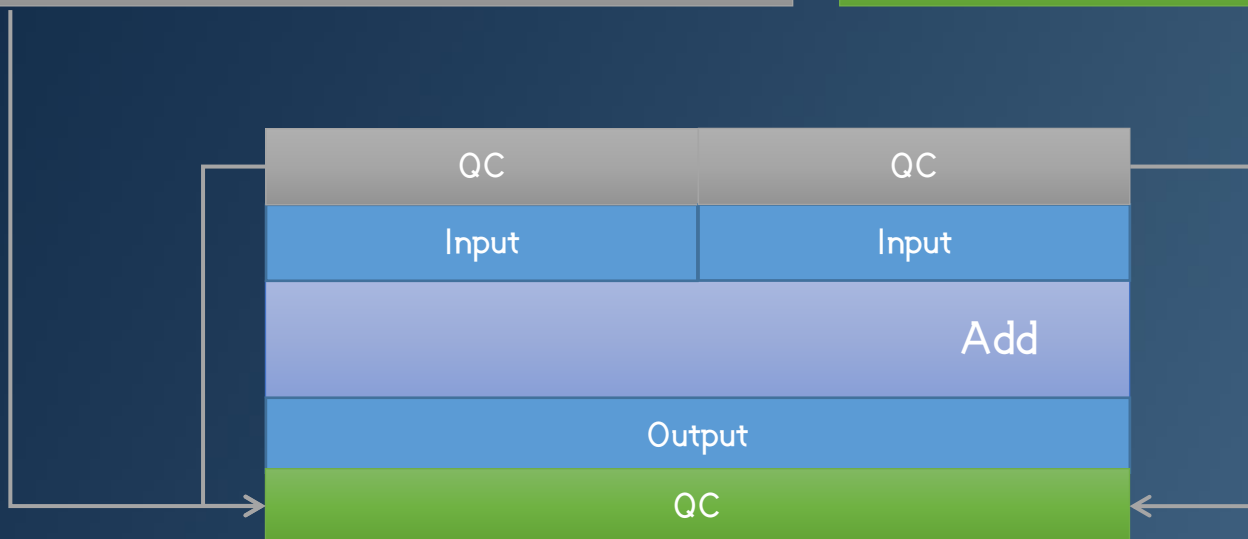

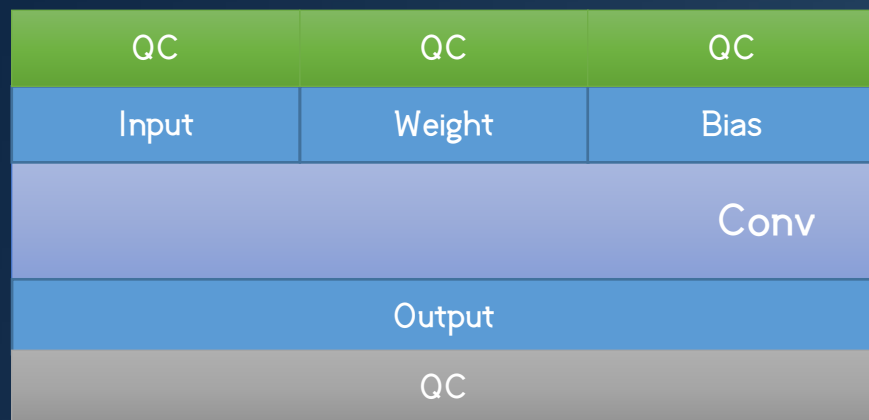
6.2.4 联合定点

Union - Quantize



6.2.4 联合定点




Union - Quantize



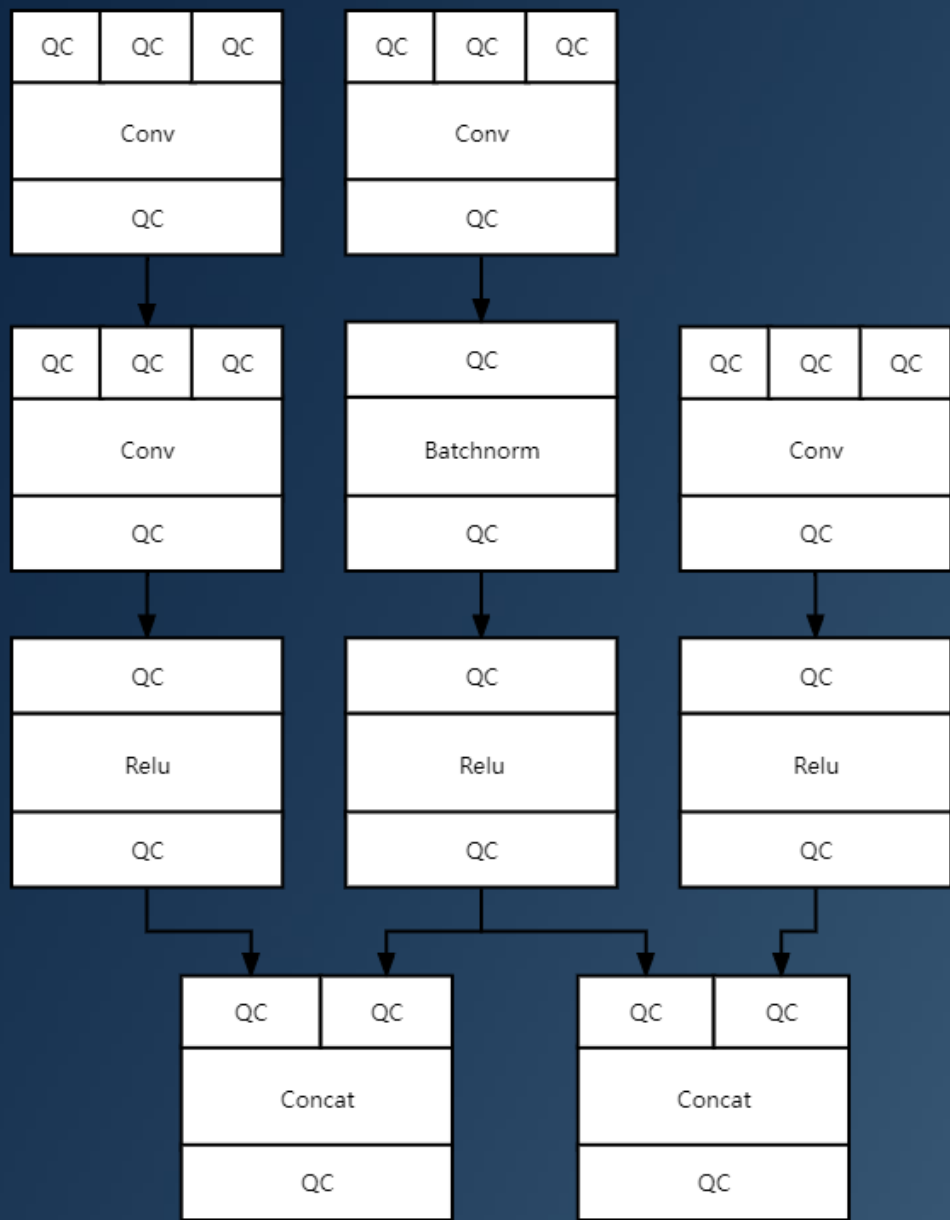
ConvAdd 联合定点 (如可能)

6.2.4 联合定点

Union - Quantize



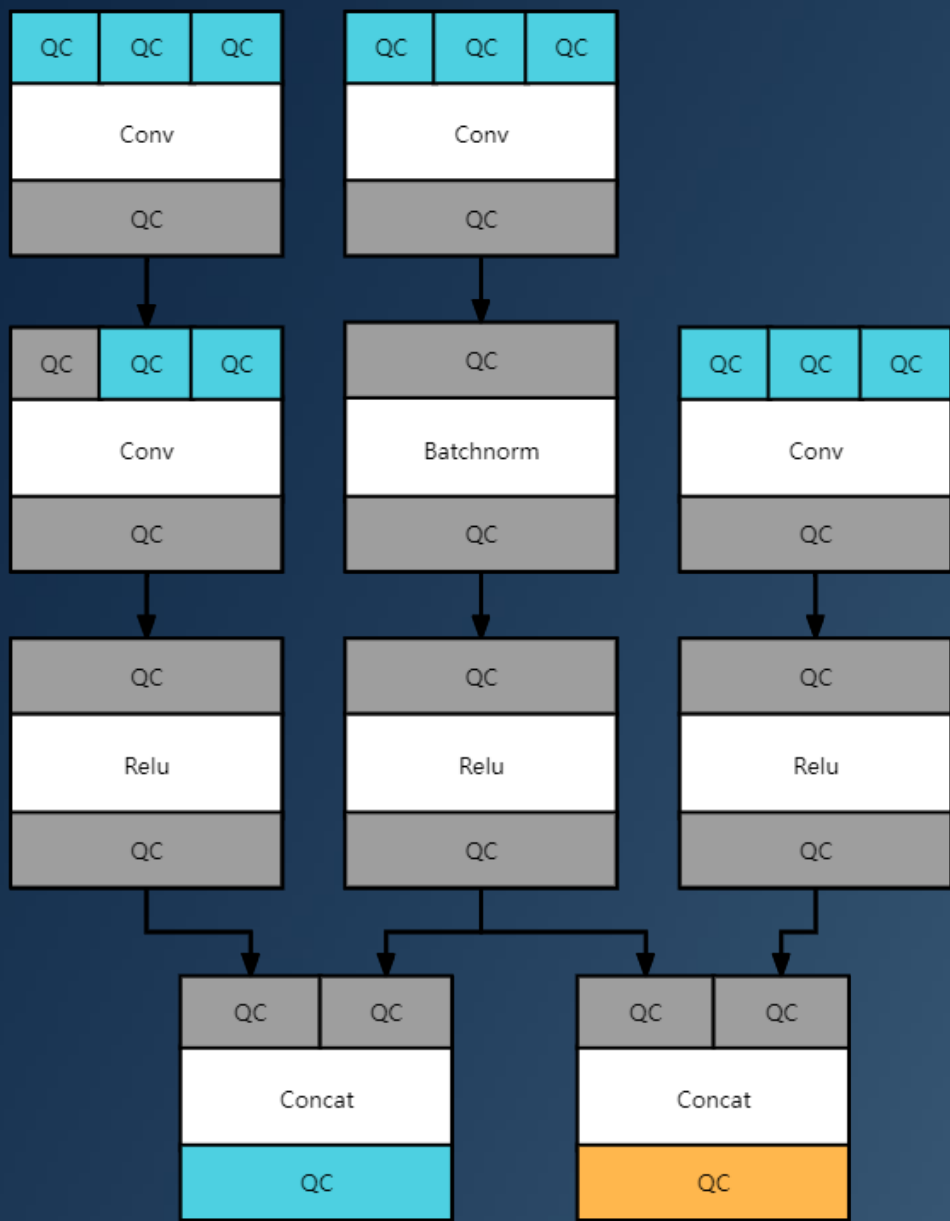
ConvAdd 联合定点 (如可能)



已知硬件存在

- conv - conv融合
- conv - relu融合
- conv - batchnorm融合
- concat联合定点

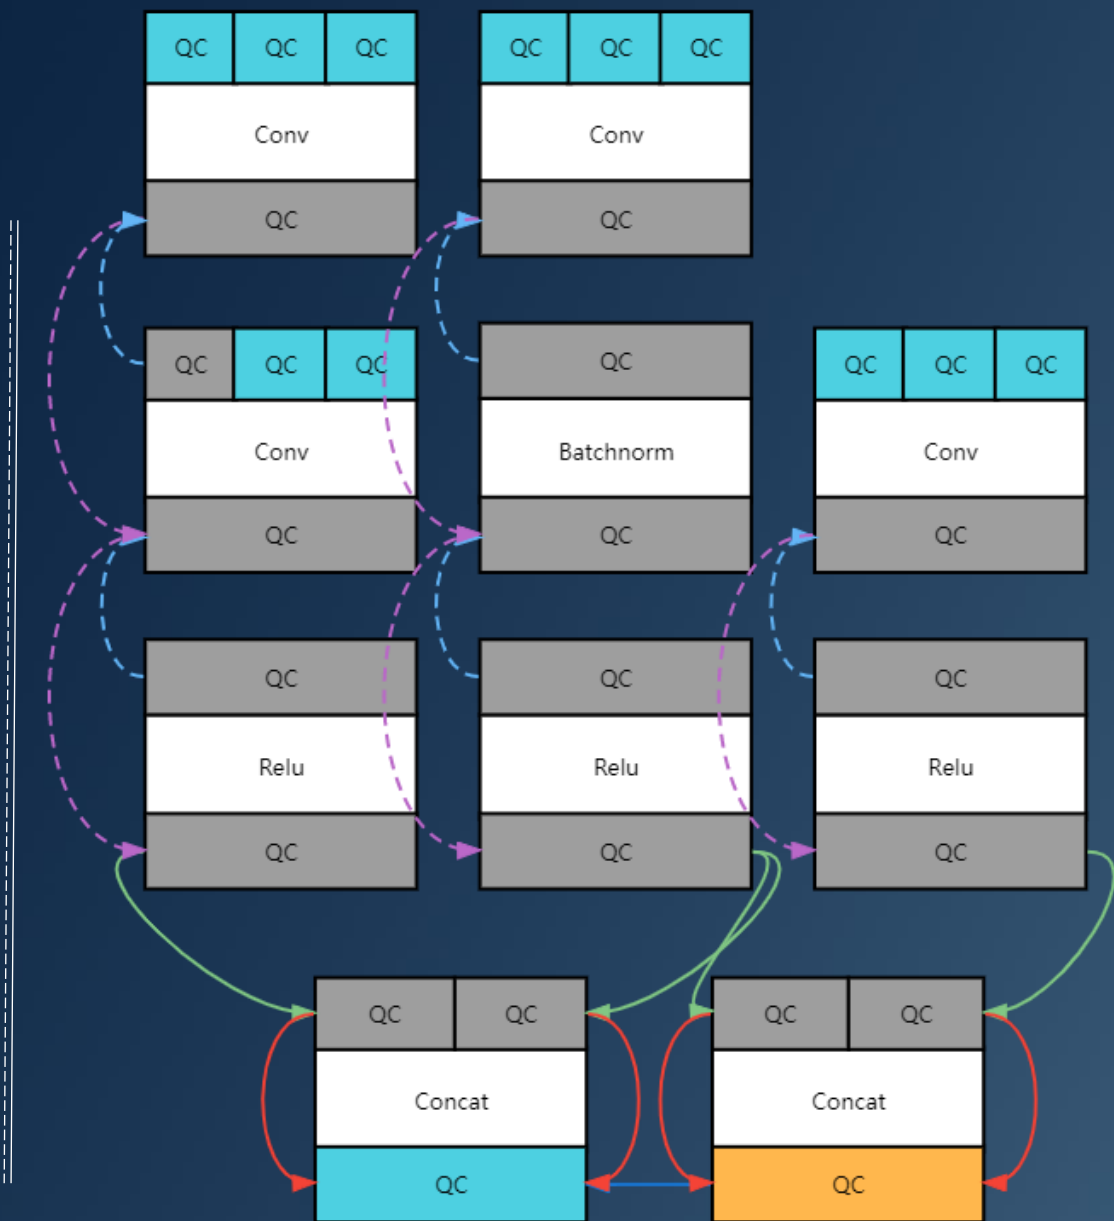
在如此规则下，左图应当如何联合定点？



已知硬件存在

- conv - conv融合
- conv - relu融合
- conv - batchnorm融合
- concat联合定点

在如此规则下，左图应当如何联合定点？




已知硬件存在

- conv - conv融合
- conv - relu融合
- conv - batchnorm融合
- concat联合定点

在如此规则下，左图应当如何联合定点？

联系我们

<https://github.com/openppl-public>



广告位招租



微信群



QQ群 (入群密令OpenPPL)